

Architectural Design

Balwinder Sodhi

MASTERING SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers



Overview of Major Architectural Styles

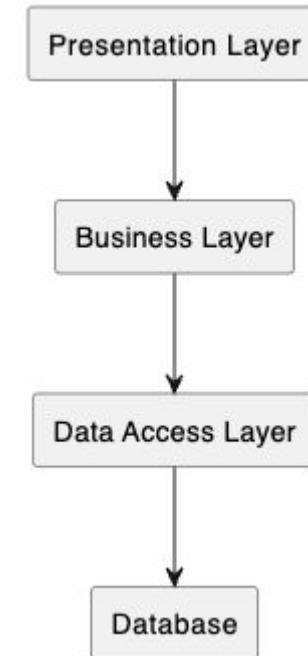
Balwinder Sodhi

Layered (N-Tier) Architecture

MASERING SOFTWARE DEVELOPMENT

A Guide For Early Career Engineers

- Key Ideas
 - System divided into layers with distinct responsibilities.
 - Common layers: Presentation, Application/Domain, Data Access.
 - Strict or relaxed layering depending on allowed dependencies.
- Why It's Used
 - Clean separation of concerns.
 - Easier maintainability and testing.
 - Ideal for CRUD + business logic apps.
- Typical Pitfalls
 - Layers become anemic (pure pass-through).
 - Too much coupling when layers leak details.
 - Inefficient for high-performance paths.



Client-Server Architecture

Key Ideas

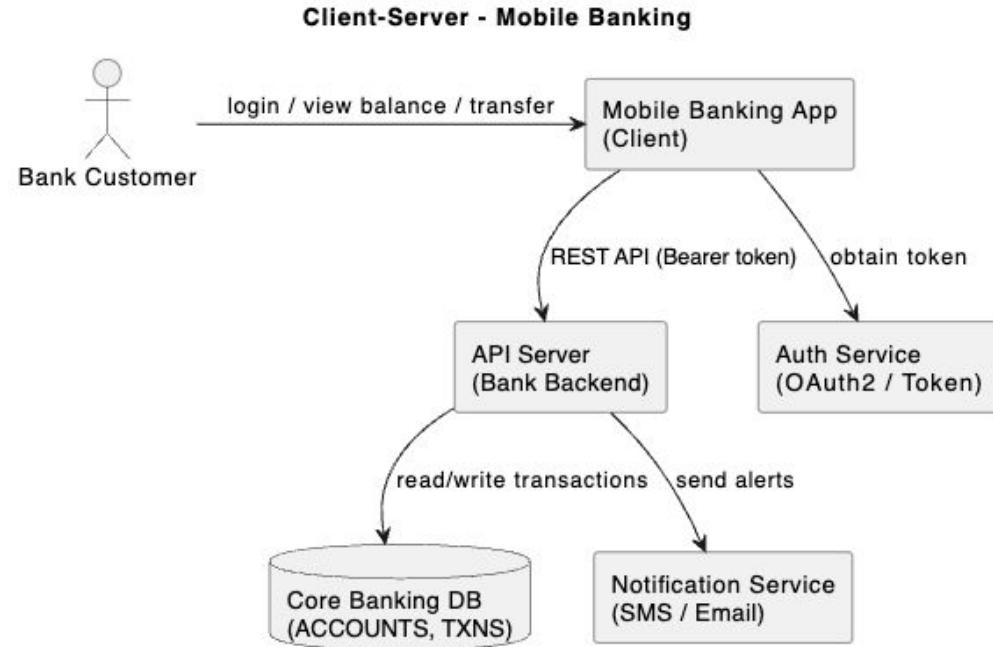
- Clients request services; servers provide them.
- General pattern behind web, distributed, networked apps.

Strengths

- Centralized control and data management.
- Works for synchronous and asynchronous interactions.

Pitfalls

- Server becomes bottleneck.
- Requires scaling strategies: horizontal, caching, load-balancing.



Microservices Architecture

● Key Ideas

- Independent services communicating over lightweight protocols.
- Autonomous deployment and scaling.

● Strengths

- High modularity.
- Polyglot freedom (services choose tech stack).
- Independent scaling and fault isolation.

● Pitfalls

- Distributed system complexity.
- Network failures, versioning, data consistency issues.
- DevOps maturity required.

Event-Driven Architecture

● Key Ideas

- Work triggered by events; decoupled producer–consumer patterns.
- Supports asynchronous, scalable workflows.

● Strengths

- Loose coupling.
- High scalability and resilience.

● Pitfalls

- Harder debugging and tracing.
- Eventual consistency issues.

Balwinder Sodhi

Service-Oriented Architecture (SOA)

● Key Ideas

- Can be thought of as a predecessor of *microservices* — kind of *macroservices*
- Enterprise services with governance, standardization (SOAP, WS-*).

● Strengths

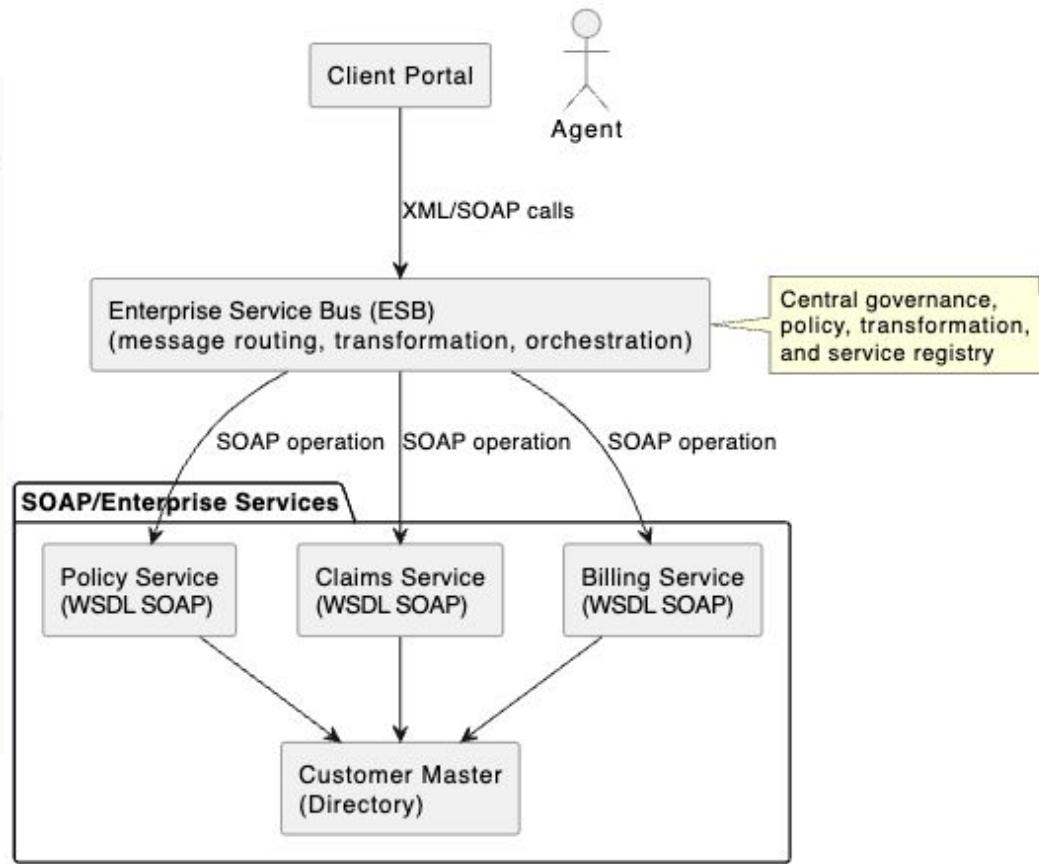
- Interoperability across enterprise systems.
- Strong contract enforcement.

● Pitfalls

- Heavy tooling.
- XML/SOAP overhead.

Balwinder Sodhi

SOA - Enterprise Insurance (ESB style)



Monolithic Architecture

MASERING SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers

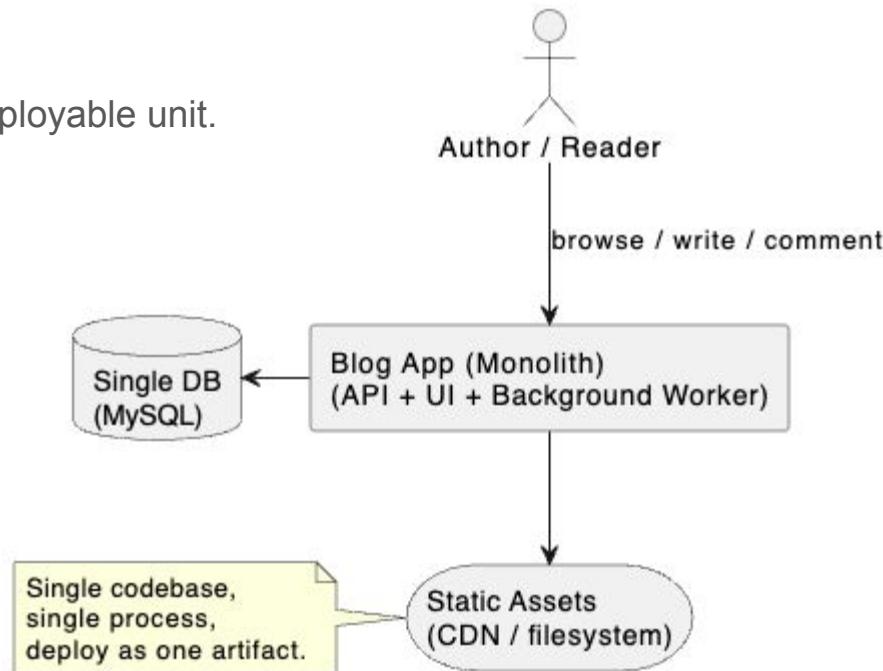
- Key Ideas
 - All functionality deployed as a single deployable unit.

- Strengths
 - Simple development and deployment.
 - Lower operational overhead.

- Pitfalls
 - Slower deployments as app grows.
 - Hard to isolate failures.
 - Scalability is “whole app only”.

Balwinder Sodhi

Monolithic - MVP Blog Platform



Component-Based Architecture

- Key Ideas

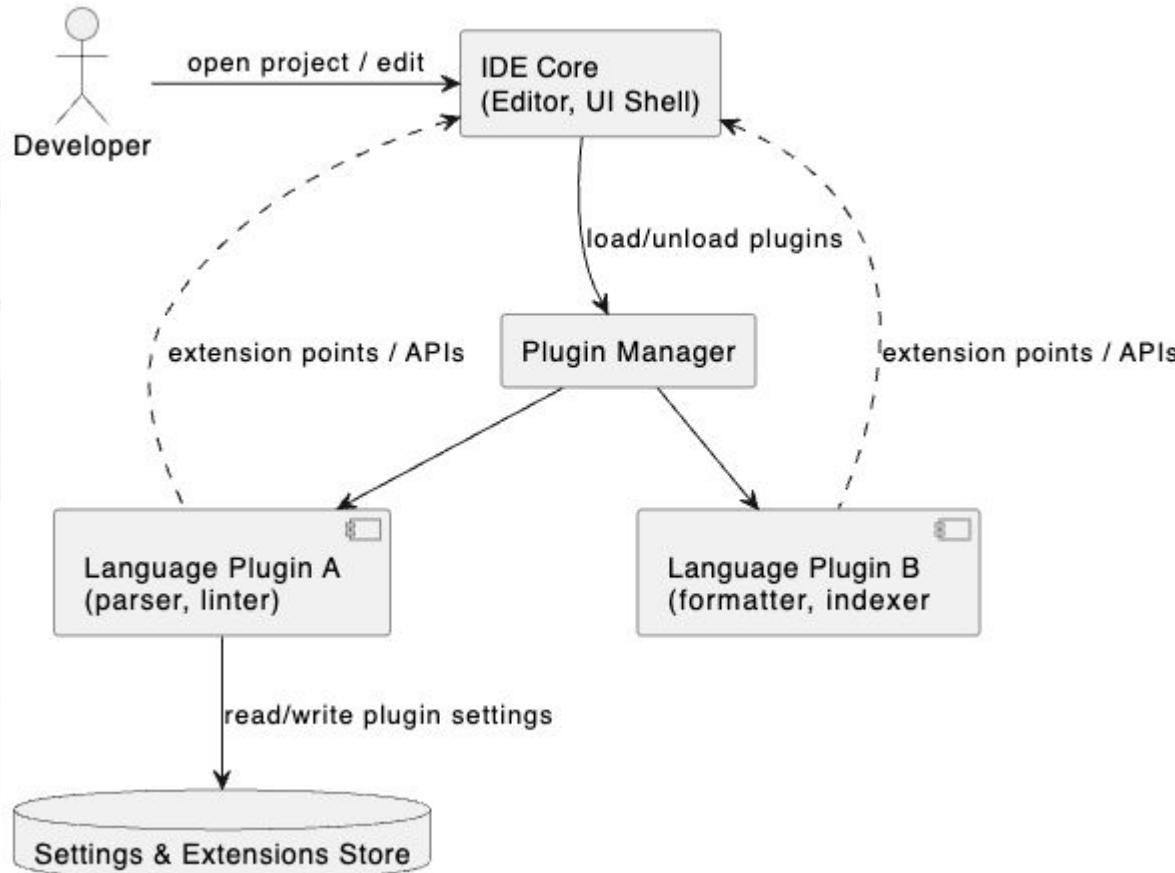
- System composed of reusable, replaceable components.
- Interfaces define the component contracts.

- Use Cases

- UI systems, plug-in systems, domain componentization.

Balwinder Sodhi

Component-Based - IDE with Plugins



Peer-to-Peer (P2P) Architecture

● Key Ideas

- Nodes act as both clients and servers.
- Resource sharing without a central authority.

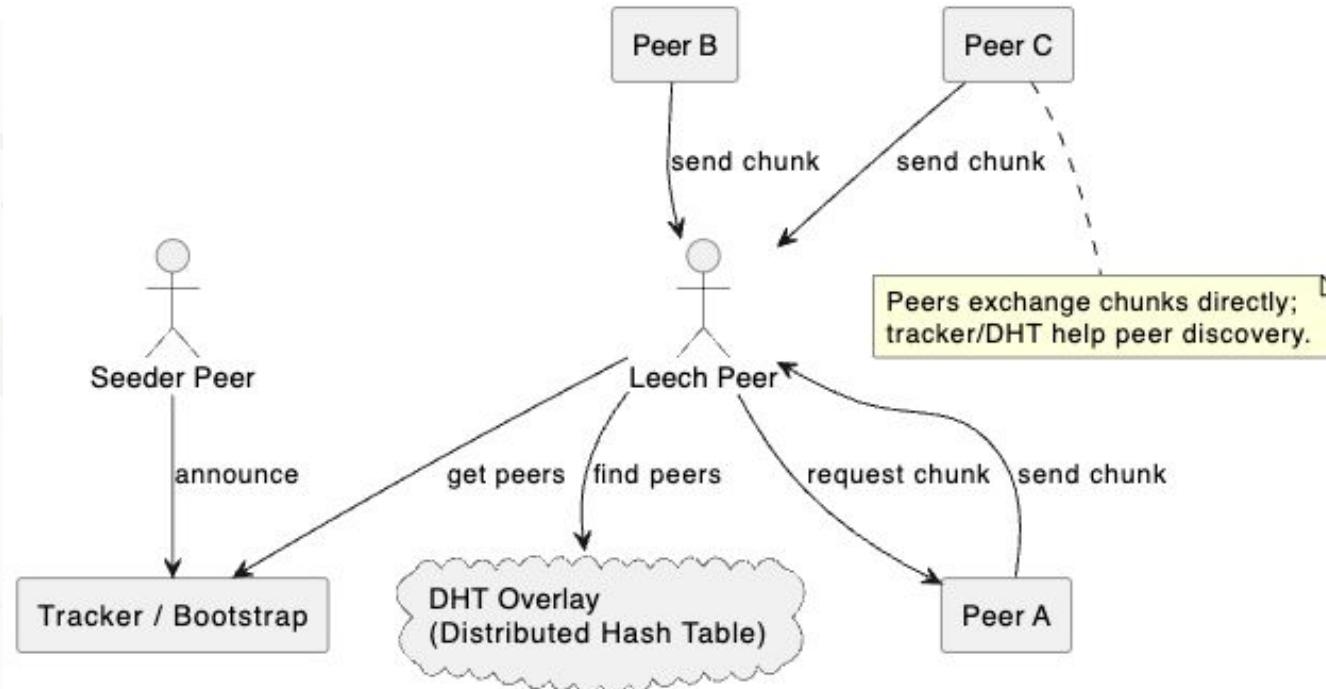
● Use Cases

- File sharing networks.
- Distributed ledgers.
- Decentralized collaboration tools.

[Go to in-depth details](#)

Balwinder Sodhi

Peer-to-Peer - File Sharing Swarm



Repository / Blackboard Architecture

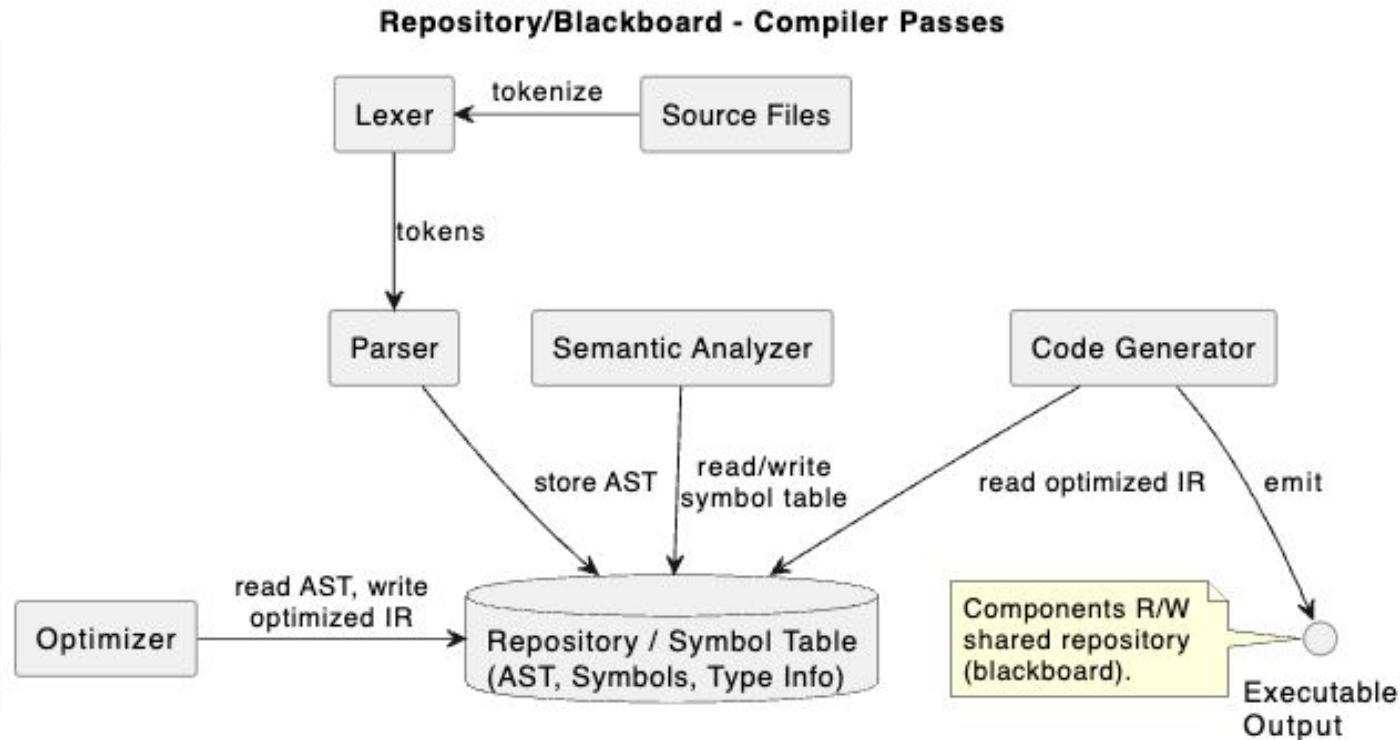
● Key Ideas

- Shared knowledge base accessed by multiple independent components.
- Used when you need incremental, cooperative problem solving.

● Examples

- Compilers (symbol table + passes).
- AI expert systems.

Balwinder Sodhi



Pipeline / Data Flow Architecture

- Key Ideas

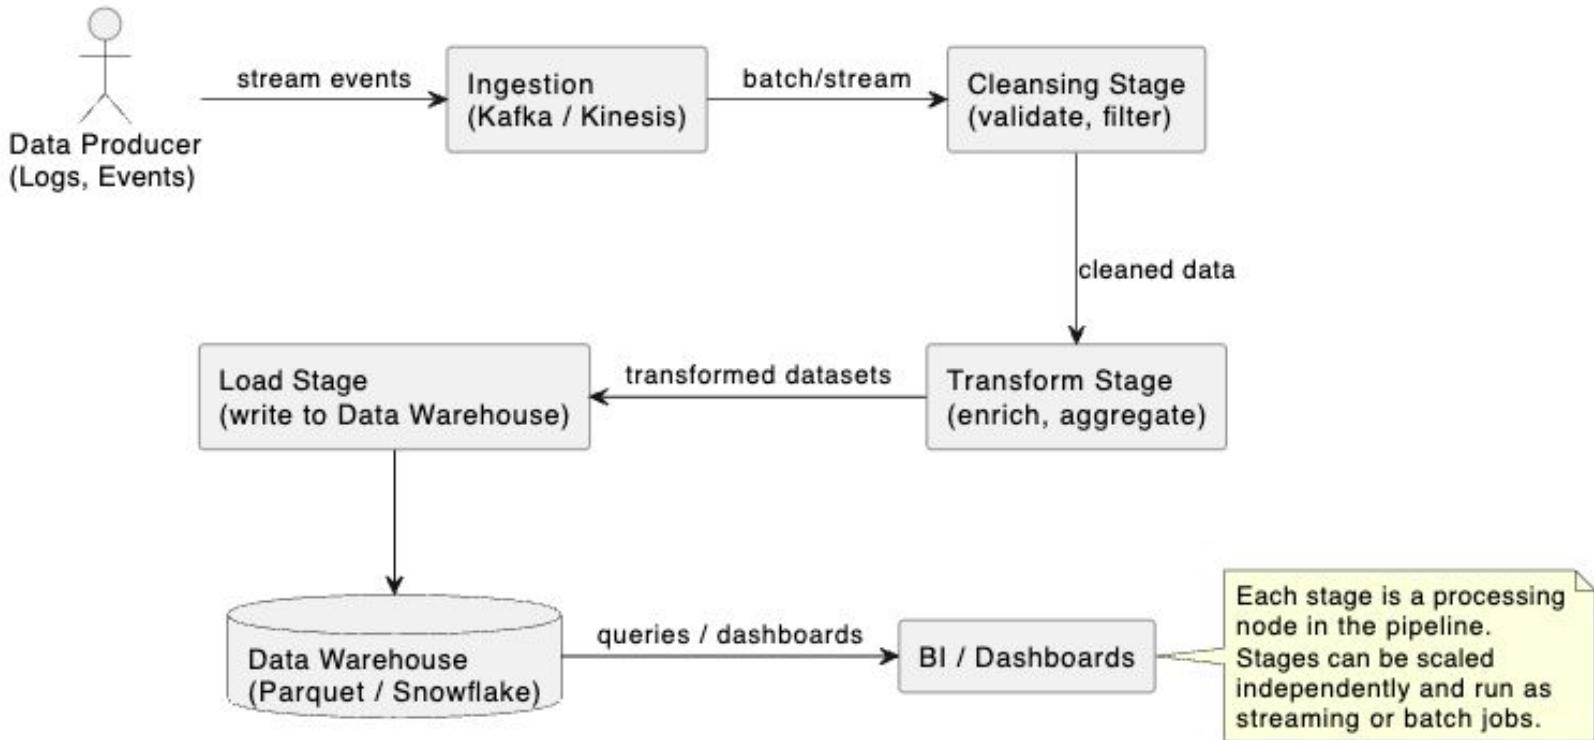
- Data flows through a series of processing stages.
- Each stage transforms data and passes to next.

- Use Cases

- ETL systems, compilers, ML pipelines.

Balwinder Sodhi

Pipeline/Data-Flow - ETL Analytics



MASTERING SOFTWARE
DEVELOPMENT

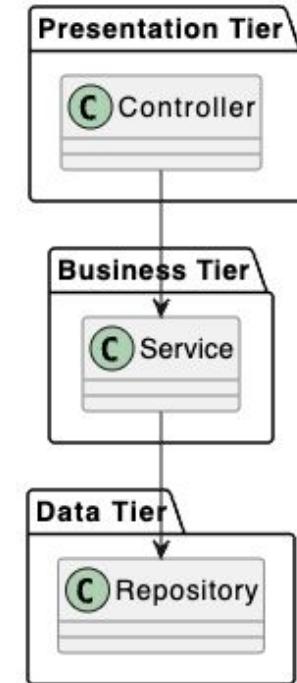
A Guide For Early Career Engineers

Selected Architectural Styles in Depth

Balwinder Sodhi

Multi-Tier (N-Tier) Architecture

- Often seen as “simple,” but:
 - the qualities they provide come from certain structural rules
 - these rules directly influence runtime behavior and organizational workflows
- How the Characteristics Emerge
 - Separation of concerns through enforced boundaries
 - Each tier has a single responsibility (UI, domain logic, data operations).
 - Enforced via directory structures, interfaces, or physical deployment boundaries.
 - It reduces accidental coupling because developers cannot easily “reach around” layers.
 - ... continued



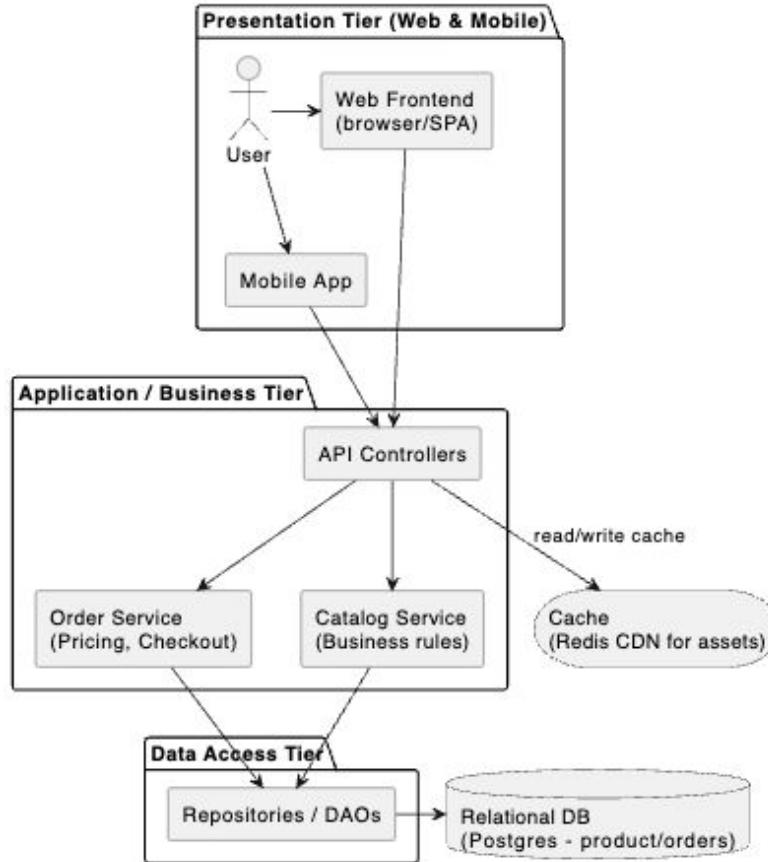
MASTERING SOFTWARE DEVELOPMENT

A Guide For Early Career Engineers



Balwinder Sodhi

Layered (N-Tier) - E-commerce Web App



Multi-Tier (N-Tier) Architecture

- How the Characteristics Emerge

- Controlled data flow

- All data typically flows top → down for reads and writes.

- This gives predictable call chains and simplifies debugging.

- Stateless middle tiers

- Business logic often designed stateless so instances can be replicated behind load balancers.

- Statelessness → horizontal scaling emerges almost naturally.

- Shared DB provides consistency

- A single relational DB ensures strong consistency for most workloads.

- The architecture itself encourages ACID-style data management.

Multi-Tier (N-Tier) Architecture: How Quality Attributes Are Achieved

Maintainability

- Because dependencies flow downward, changes in lower tiers don't break upward layers.
- Replaceability of layers (e.g., moving from SQL to NoSQL) becomes possible if contracts are stable.

Security

- Sensitive components (DB, secrets) reside in inner network zones.
- Network policies create “security rings” around critical resources.

Performance

- Caching can be placed strategically:
 - UI cache (CDN)
 - Application-tier local cache
 - Database read replicas

Availability

- Stateless tiers let you scale out horizontally.
- With load balancers, you gain resilience against instance failures.

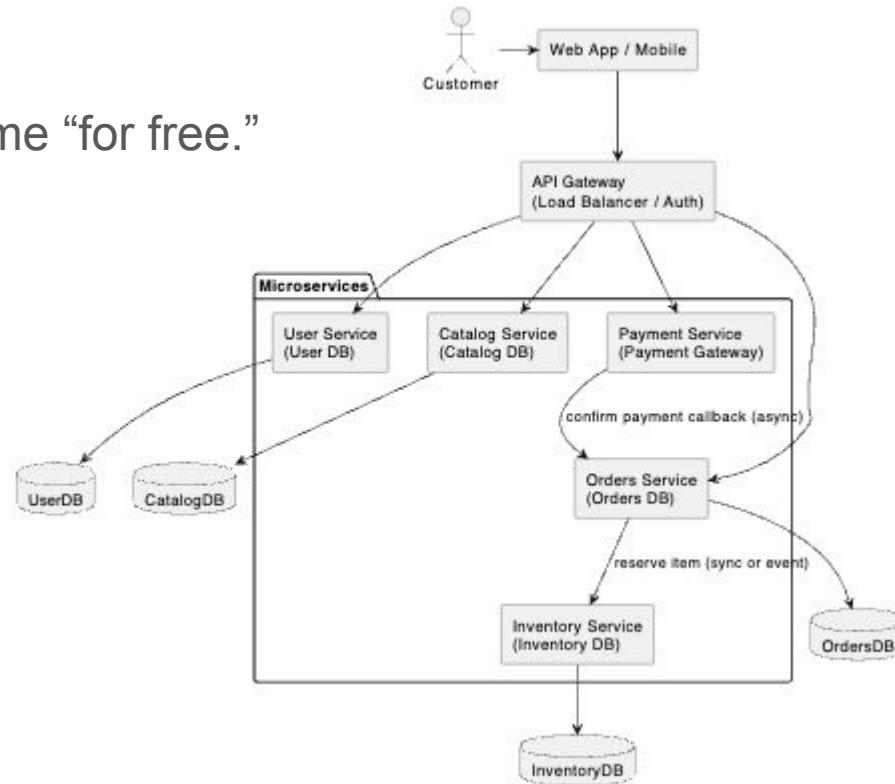


Microservices Architecture

- MAS
TER
DEVELOPMENT
A Guide For Early Career Engineers
- Microservices qualities do not come “for free.”
They arise from:

- bounded contexts
- independent deployability
- distributed data, and
- asynchronous communication.

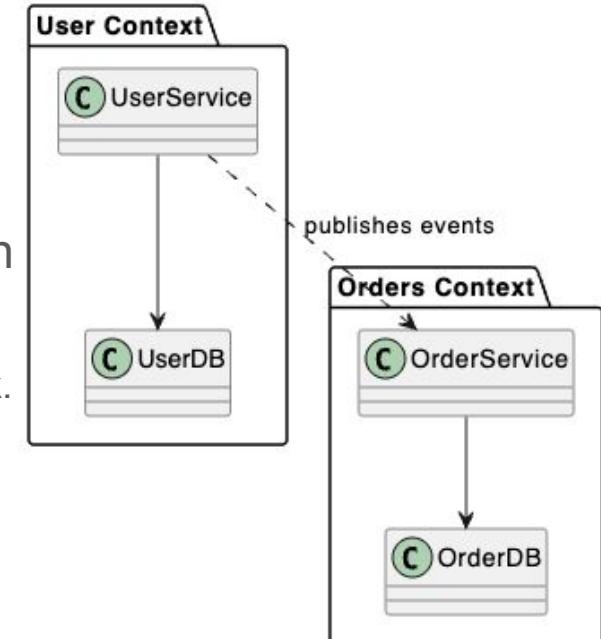
Microservices - Online Retailer



Balwinder Sodhi

Microservices – How the Characteristics Emerge

- Bounded contexts enforce modularity
 - Each service owns a clear, well-defined domain boundary.
 - This prevents business logic bloat in shared modules.
 - Teams work independently without merging conflicts.
- Independent deployment through process isolation
 - Each service is its own deployable unit.
 - CI/CD per service → agility and parallel development.
 - Rolling updates and canary releases reduce downtime risk.
- Distributed data autonomy
 - Each service owns its storage (polyglot persistence).
 - Eliminates shared DB bottlenecks and schema conflicts.
- Communication via APIs and events
 - Services interact through REST/gRPC for sync calls.
 - Events reduce direct coupling and support eventual consistency.



Microservices – How Quality Attributes Are Achieved

Scalability

- Services that experience high load (e.g., checkout) can be scaled independently.
- Hot paths can use faster protocols (gRPC, binary).

Resilience

- Failures stay contained:
 - Bulkheads (resource isolation)
 - Circuit breakers
 - Retry/backoff patterns
- Event-driven communication prevents cascading failures.

Modifiability

- A service can switch frameworks, languages, or storage without impacting others.
- Backward-compatible APIs allow safe evolution over time.

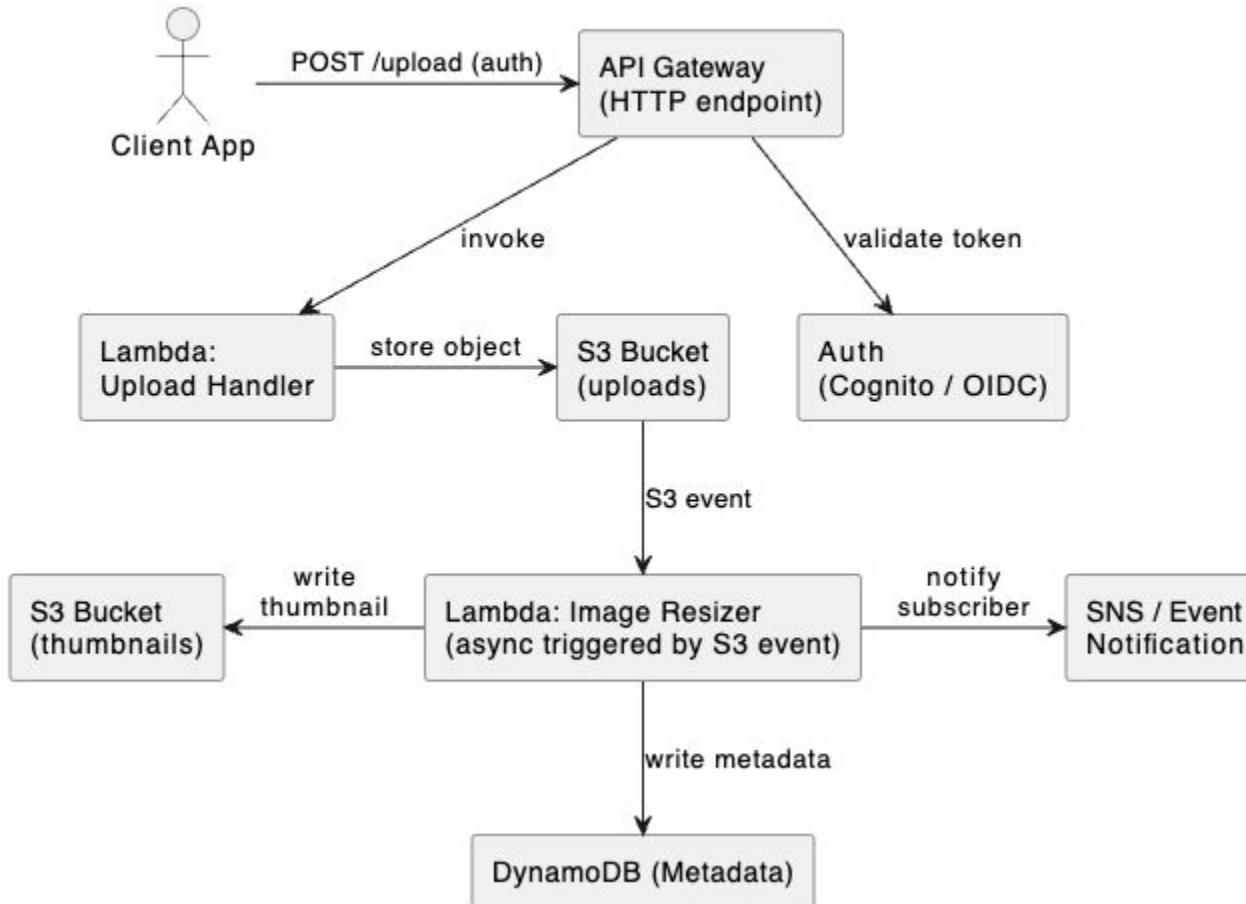
Deployability

- Small size → fast build and deploy times.
- Per-service pipelines eliminate global coordination.

Serverless Architecture

- Serverless qualities arise from:
 - event-driven compute
 - managed runtime environments, and
 - auto-scaling provided by the cloud provider.
- Functions triggered by events
 - Code runs only when invoked by HTTP, queue, cron, or cloud events.
 - Removes the need for dedicated servers or long-lived processes.
- Auto-scaling through ephemeral functions
 - Cloud provider spins up more function instances when traffic spikes.
 - No manual scaling logic to implement.

Serverless - Image Processing / Thumbnailing



Serverless Architecture

- Fine-grained billing model
 - You pay only for execution time and resource consumption.
 - Idle time costs zero.
- Shared responsibility model
 - Provider handles:
 - OS patching
 - Server provisioning
 - Load balancing
 - Basic security hardening

Balwinder Sodhi

Serverless Architecture – How Quality Attributes Are Achieved

Scalability

- Functions scale horizontally without developer involvement.
- The platform auto-provisions instances.

Cost Efficiency

- Short-lived operations avoid overprovisioning.
- Great for bursty or erratic workloads (IoT events, batch jobs).

Availability

- Cloud providers run functions across multiple AZs.
- No single-instance failures.

Time-to-Market

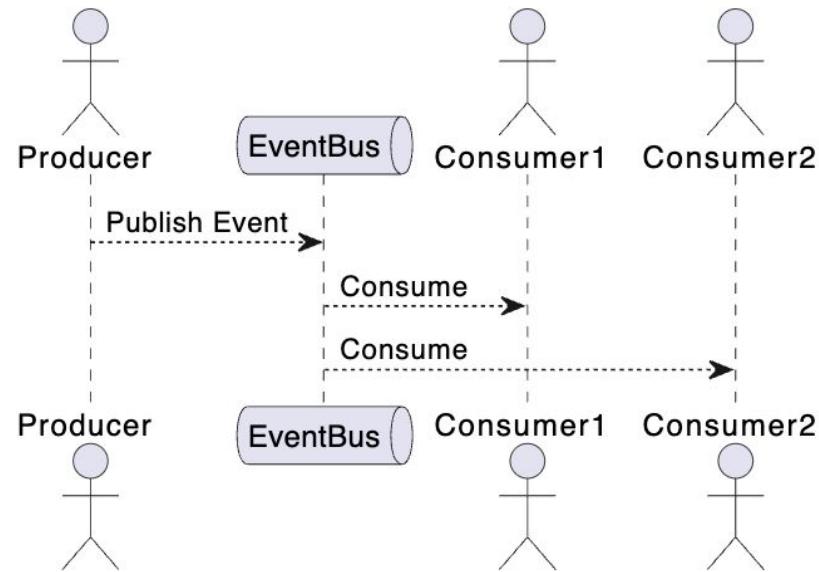
- Teams focus solely on business logic.
- No need for container, VM, or infrastructure management.

Event-Driven Architecture

- How the Characteristics Emerge

- Producers unaware of consumers
 - Producers only publish events; they don't know who consumes them.
 - New consumers can be added without modifying producers.
- Event logs as the backbone
 - Systems like Kafka provide:
 - Ordered partitions
 - Durability
 - Replayability

... contd.



Event-Driven: Characteristics – How they emerge

- Scaling via consumer groups
 - Multiple consumers in a group share workload.
 - Horizontal scaling becomes a configuration issue.
- Eventual consistency
 - Instead of synchronous global transactions, services update their own state after events.

Balwinder Sodhi

Event-Driven: Quality Attributes

- **Loose coupling**

- Changes to consumer logic never require touching producers.
- Encourages autonomous teams.

- **Scalability**

- Event brokers handle massive throughput.
- Consumer groups allow elastic scaling.

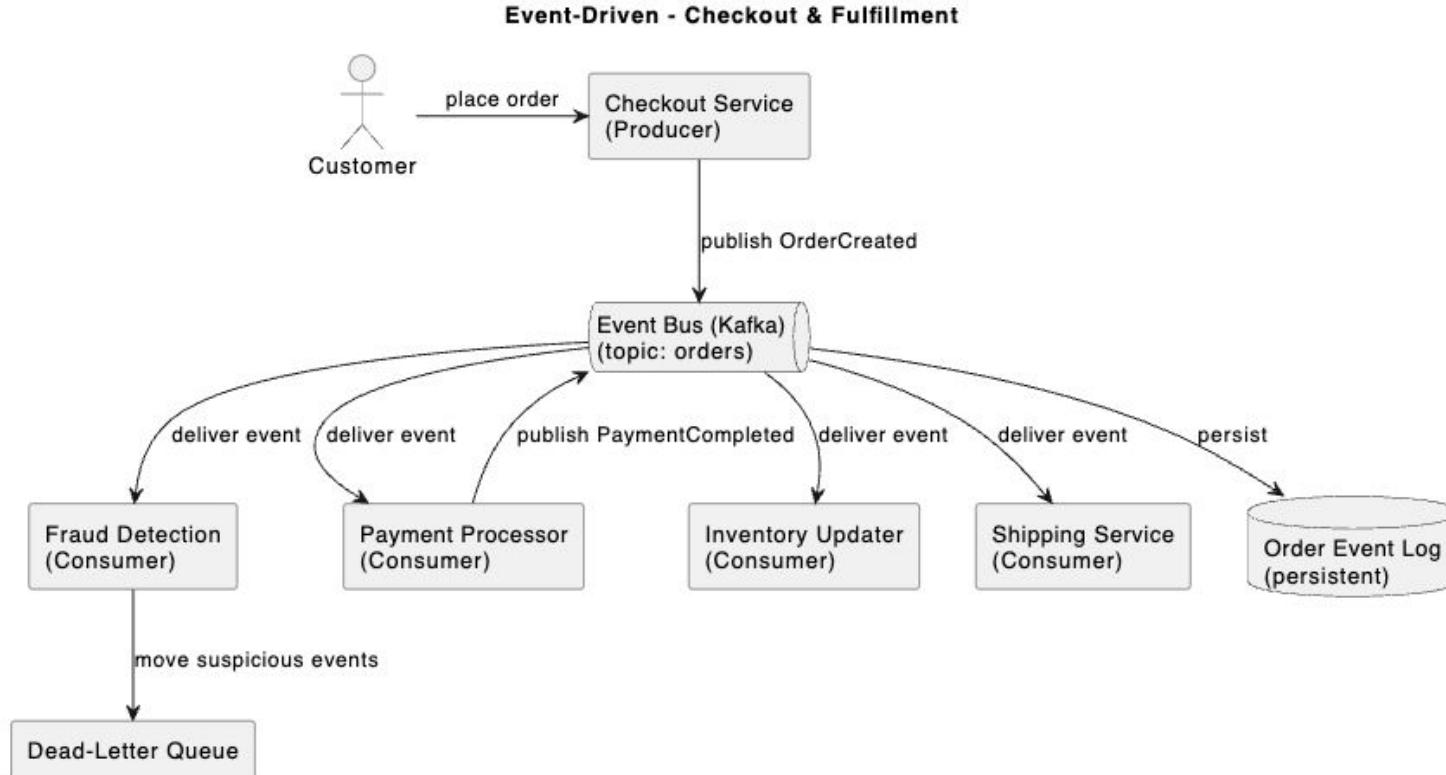
- **Resilience**

- Event logs preserve events even during outages.
- Consumers can resume from last offsets.

- **Extensibility**

- Adding new downstream processing is straightforward:
 - Real-time analytics
 - Monitoring
 - Fraud detection
 - Alerts

Event-driven Architecture Example



Peer-to-Peer Architecture

- P2P systems rely on:
 - decentralization
 - distributed state, and
 - peer routing protocols.

We have seen the basic details of P2P [earlier](#)

Balwinder Sodhi

P2P: How the Characteristics Emerge

- No central coordinator
 - Peers discover each other using bootstrapping nodes or gossip protocols.
 - Responsibilities are distributed (storage, computation, routing).
- Data replication among peers
 - Replication ensures high availability without servers.
 - Missing data can be reconstructed from neighbor peers.
- Overlay networks (e.g., DHTs)
 - Distributed Hash Tables enable:
 - $O(\log N)$ lookup
 - Even distribution of keys
- Self-scaling
 - More peers = more capacity.
 - Aggregate bandwidth and storage grow automatically.

P2P: How Quality Attributes Are Achieved

- **Resilience**

- No single point of failure because responsibilities are decentralized.

- **Availability**

- Robust replication ensures data survivability.

- **Scalability**

- As more peers join, the network gains more resources.
 - Non-linear scaling for file-sharing (BitTorrent swarms).

- **Decentralization**

- Useful in censorship-resistant or trust-minimized systems.

Balwinder Sodhi

Choosing the Right Architecture

- Key Decision Factors

- Functional requirements
- Non-functional goals (latency, throughput, availability)
- Team skill & organization
- Operational maturity (DevOps, observability)
- Regulatory constraints
- Expected scale & workload shape

- Key Trade-offs

- Monolith ↔ Microservices
- Event-Driven ↔ Request/Response
- Serverless ↔ Containerized
- Strong consistency ↔ Eventual consistency

MASTERING SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers



Cross-Cutting Architecture Issues

Balwinder Sodhi

Data Access Abstraction

MASTER SOFTWARE DEVELOPMENT

A Guide For Early Career Engineers

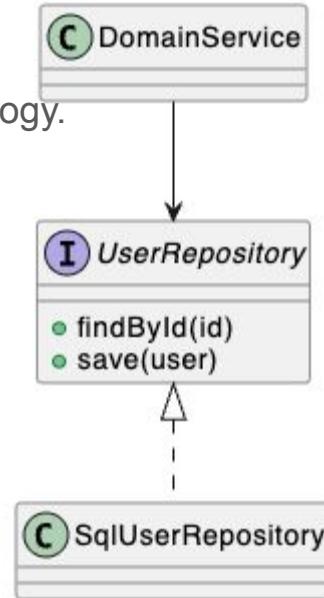
Balwinder Sodhi

- Why

- Reduce coupling between business logic and database schema/technology.
- Enable database migrations or vendor switches.
- Improve testability by mocking data access.

- How

- Repository or DAO pattern hides SQL/ORM specifics.
- Use interfaces to abstract DB operations.
- Use mappers to convert between domain and persistence models.



Authentication and Authorization

MASER SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers

- Why
 - Identify users and enforce permissions.
 - Required for security, auditability, compliance.

- How
 - OAuth2/OIDC: token-based workflows with identity providers.
 - Access tokens signed using JWTs.
 - Role-based access control (RBAC) via roles → permissions.
 - Attribute-based access control (ABAC) using runtime policies.

- Operational Considerations
 - Token rotation and expiry.
 - Store passwords using slow hashing (bcrypt, Argon2).
 - Enforce MFA, rate-limiting login attempts.

Securing Data in Transit and At Rest

MASTER SOFTWARE DEVELOPMENT

A Guide For Early Career Engineers

BY Balwinder Sodhi

● Why

- Prevent leakage through network sniffing or disk compromise.
- Regulatory mandates (GDPR, HIPAA, PCI DSS).

● How

- TLS 1.2+ for transport; enforce HSTS.
- Encrypt sensitive fields with AEAD ciphers (AES-GCM).
- Store keys using KMS or HSM for rotation and access control.
- Encrypted backups and secure snapshot storage.

Balwinder Sodhi

Input Validation & Sanitization

MASER SOFTWARE DEVELOPMENT

A Guide For Early Career Engineers

Why

- Prevent injection attacks.
- Ensure system receives data in predictable formats.
- Improves usability

How

- Strict schema validation at trust boundaries.
- Escape untrusted input (HTML escape, SQL bind parameters).
- Treat all external input as hostile.
- Use allowlists instead of blocklists.

Balwinder Sodhi

Logging and Auditing

MASER SOFTWARE
DEVELOPMENT

- Why
 - Observability, security, compliance.
 - Helps troubleshooting and forensic analysis.

- How

- Use structured logs (JSON).
- Include correlation IDs to trace multi-service flows.
- Sensitive-information filters to prevent leaking secrets.
- Centralized log aggregation via ELK, Loki, or CloudWatch.

Balwinder Sodhi

Instrumentation and Metrics

MASTER SOFTWARE DEVELOPMENT

A Guide For Early Career Engineers

Why

- Lets teams monitor system health and performance.
- Detect anomalies early.

How

- Use RED (Rate, Errors, Duration) for microservices.
- Use USE (Utilization, Saturation, Errors) for infrastructure.
- Export metrics using OpenTelemetry.
- Add distributed tracing using spans across service boundaries.

Balwinder Sodhi

Concurrency and Parallelism

MASTER SOFTWARE DEVELOPMENT

A Guide For Early Career Engineers

- Why
 - Improve performance.
 - Utilize multi-core hardware effectively.

- How
 - Employ multi-threading/processing
 - Worker pools, event-driven design, queues, etc.

- Key considerations
 - Avoid shared mutable state → race conditions.
 - Use messaging or actor model for isolation.
 - Use locks sparingly; favor lock-free structures.
 - Manage thread pools carefully to avoid exhaustion.

Error Handling

MAS TER OF SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers

● Why

- Improve reliability.
- Prevent cascading failures.
- Enhance user experience.

● How

- Central error handlers at API boundaries.
- Retries with jitter to avoid thundering herd.
- Circuit breakers for remote calls.
- Error/ dead letter queue handling for failed messages.

Balwinder Sodhi

MASTERING SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers

API Design & Integration

Balwinder Sodhi

API Styles Overview

MASER SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers

- REST
- SOAP
- GraphQL
- WebSockets
- RPC (gRPC)
- Library/SDK
- File-based
- Custom Protocols
- Serverless APIs

Pavinder Sodhi

REST: Key Ideas

- Resource-oriented (/users, /orders/123)
- Stateless HTTP interactions
- Uniform verbs: GET, POST, PUT, DELETE
- Uses standard HTTP semantics and caching
- Ideal for CRUD-type operations

Balwinder Sodhi

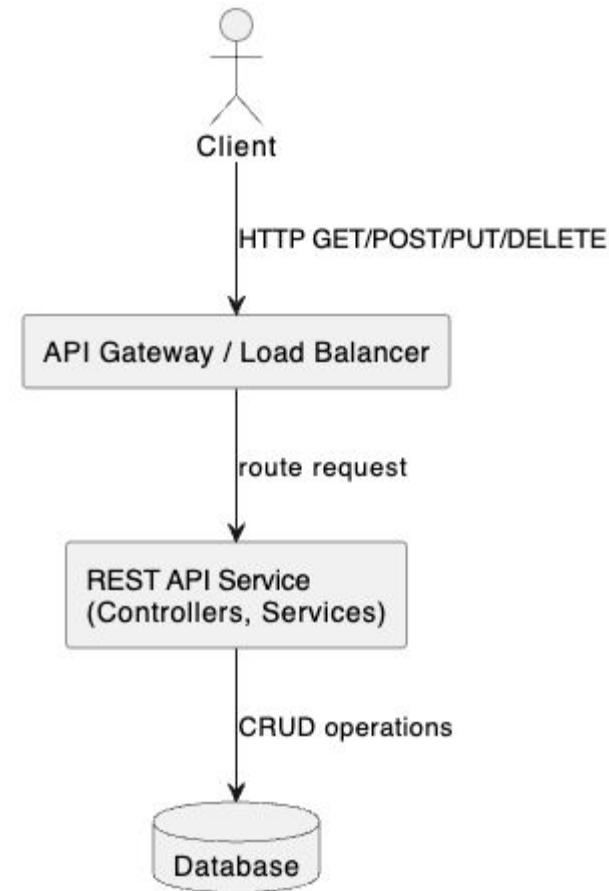
REST: When to Use

- Public APIs (web services, mobile apps)
- CRUD-heavy business systems
- When simplicity, compatibility, browser-friendliness matter
- Good default choice for most teams

Balwinder Sodhi

REST: How Typically Implemented

- HTTP server exposes resource endpoints
- JSON payloads
- Web clients or mobile apps consume via fetch/axios
- Backed by controllers → services → database
- Uses auth tokens (JWT, OAuth2)



Balwinder Sodhi

SOAP: Key Ideas

- XML-based protocol
- Strong typing via WSDL (contract-first)
- Heavy on standards: WS-Security, WS-Addressing
- Supports formal enterprise workflows

Balwinder Sodhi

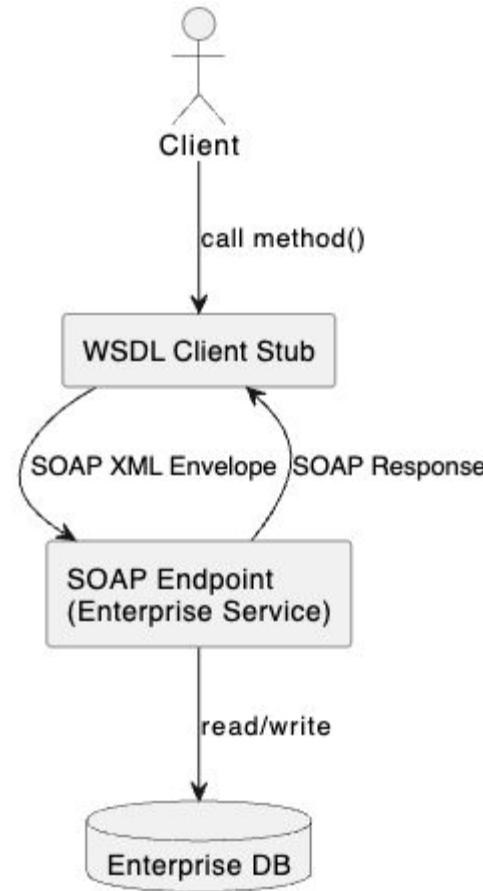
SOAP: When to Use

- Enterprises with legacy or regulated systems
- Strict schema and validation requirements
- Banking, insurance, telecom
- Systems needing guaranteed message delivery and WS-* standards

Balwinder Sodhi

SOAP: How Typically Implemented

- SOAP server exposes WSDL contract
- XML messages exchanged over HTTP or Message Queues
- Client stubs auto-generated from WSDL
- Often deployed behind enterprise ESBs



Balwinder Sodhi

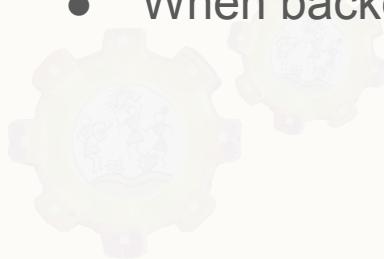
GraphQL: Key Ideas

- Client specifies exactly what data it wants
- Single endpoint (/graphql)
- Strong schema (SDL) defines types
- No over-fetching or under-fetching
- Ideal for modern UI with variable data needs

Balwinder Sodhi

GraphQL: When to Use

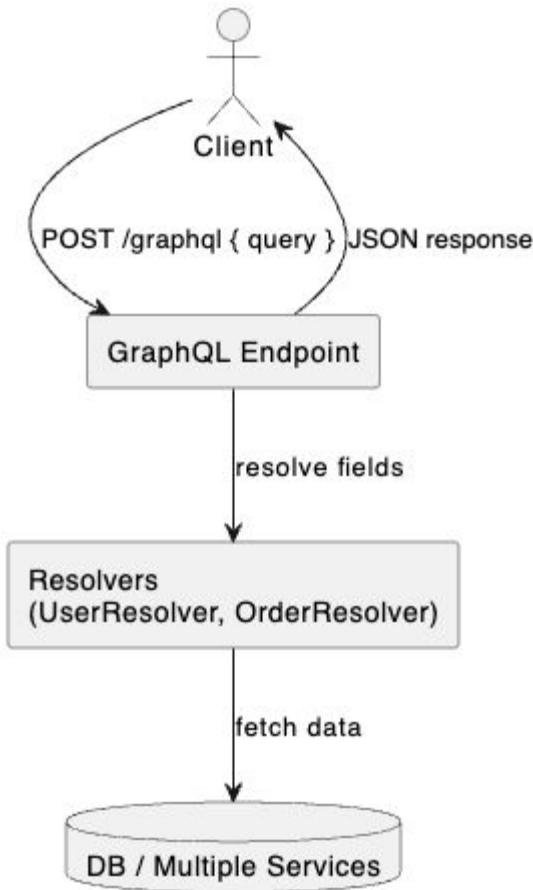
- Complex data graphs (social feeds, product catalogs)
- Mobile apps minimizing payload size
- Frontend teams want flexibility
- When backend teams want schema governance



Balwinder Sodhi

GraphQL: How Typically Implemented

- GraphQL server parses queries → resolvers
→ services → DB
- Resolvers fetch only required fields
- Supports batching, caching, pagination



Balwinder Sodhi

WebSockets: Key Ideas

- Full-duplex, persistent connection
- Real-time bidirectional communication
- No polling needed
- Ideal for continuous streams or updates

Balwinder Sodhi

WebSockets: When to Use

- Chat apps
- Multiplayer games
- Real-time dashboards
- Stock tickers, IoT telemetry

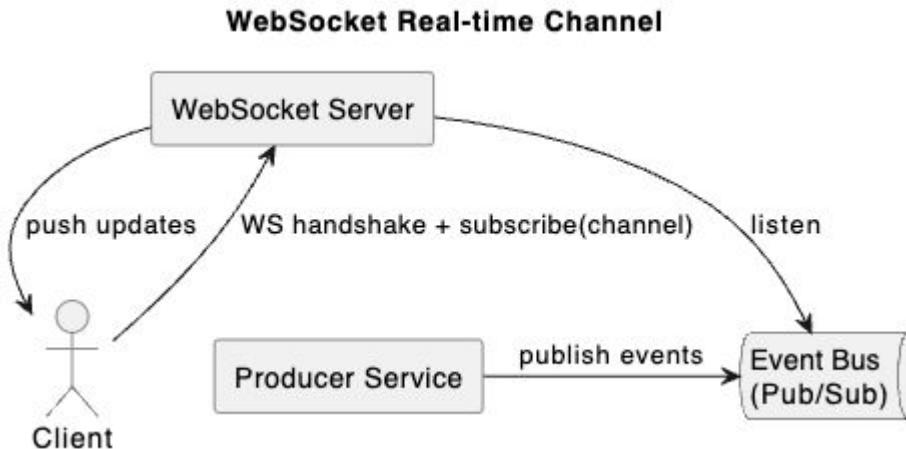
Balwinder Sodhi

WebSockets: How Typically Implemented

- Client upgrades HTTP → WebSocket handshake
- Server and client exchange messages
- Backend services push events to clients via broker or pub/sub system



Balwinder Sodhi



RPC (gRPC): Key Ideas

- Contract-first via Protocol Buffers
- Fast binary messages
- Bi-directional streaming supported
- Strongly typed stubs auto-generated

Balwinder Sodhi

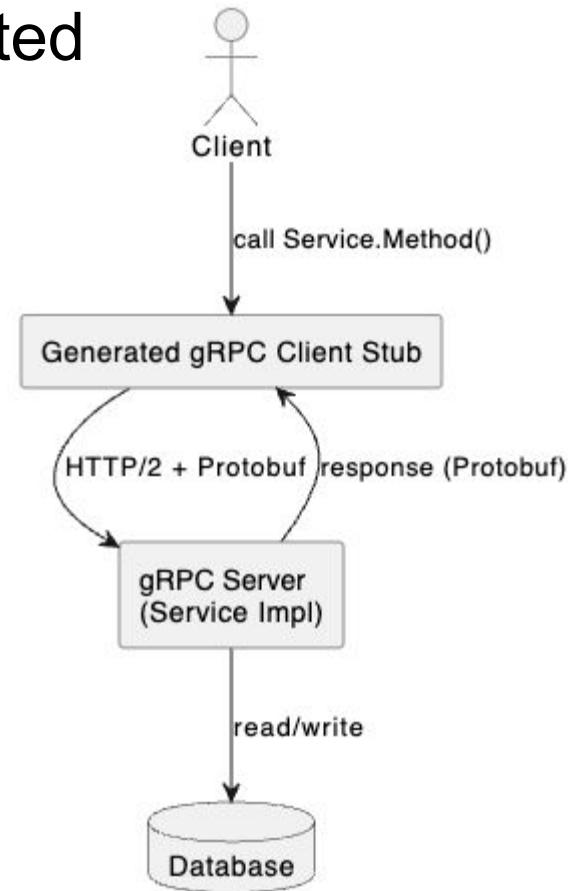
RPC (gRPC): When to Use

- High-performance inter-service calls
- Microservices backend communication
- Latency-sensitive operations
- Typed APIs between trusted services

Balwinder Sodhi

RPC (gRPC): How Typically Implemented

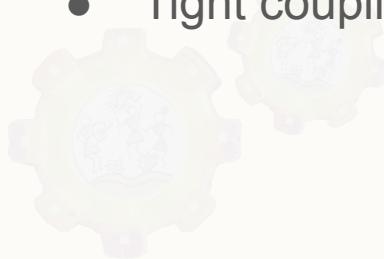
- .proto file defines request/response
- Codegen generates client/server stubs
- Server handles method calls directly
- HTTP/2 transport gives multiplexed streams



Balwinder Sodhi

Library/SDK APIs: Key Ideas

- Delivered as a code package (npm, pip, Maven, NuGet)
- Runs inside client's application runtime
- Ideal for offline logic, wrappers, utilities
- Tight coupling with client code



Library/SDK APIs: When to Use

- Domain logic as reusable modules
- Complex integrations wrapped behind easy interfaces
- Performance-sensitive local functionality (crypto, image processing)

MASER
DEVELOPMENT

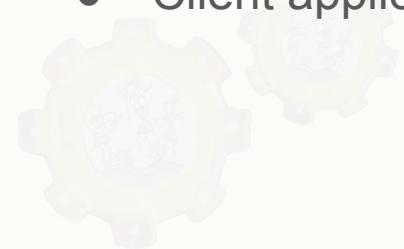
A Guide For Early Career Engineers



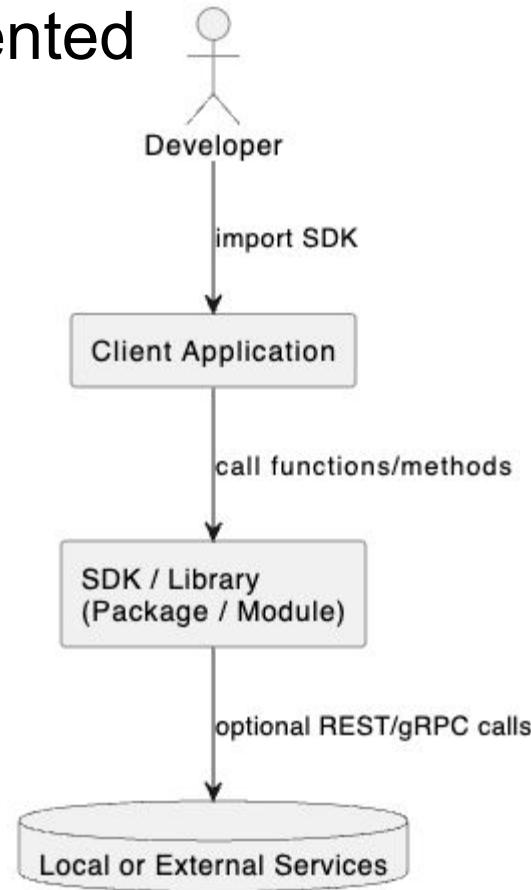
Balwinder Sodhi

Library/SDK APIs: How Typically Implemented

- MASER DEVELOPMENT
A Guide For Early Career Engineers
- Maintained as versioned library
 - Semantic versioning
 - Exposes classes, functions, interfaces
 - Client applications import and use directly



Balwinder Sodhi



File-Based APIs: Key Ideas

- Integrations via exchanged files (CSV, JSONL, XML, Parquet)
- Scheduled or ad-hoc batch transfers
- Often used for ETL or legacy integrations



Balwinder Sodhi

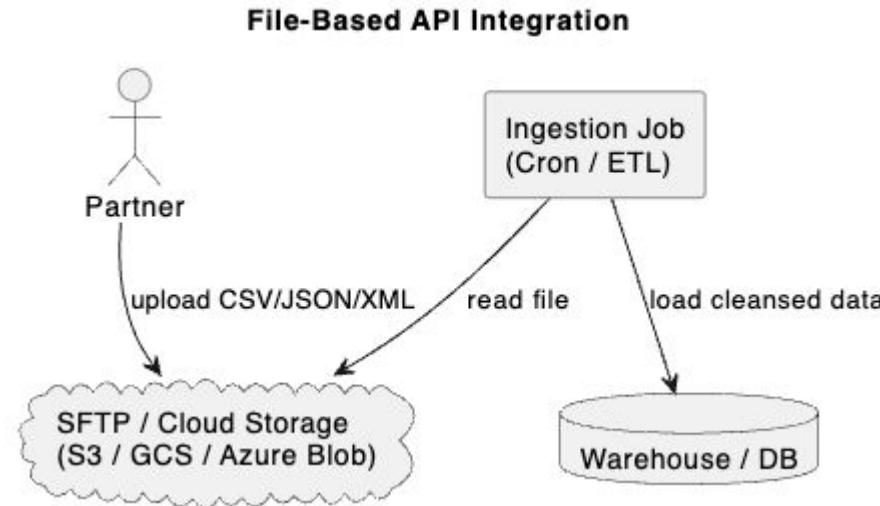
File-Based APIs: When to Use

- Bulk data exchanges
- Partners who cannot expose HTTP APIs
- BI/analytics ingestion
- Cross-company data sharing

Balwinder Sodhi

File-Based APIs: How Typically Implemented

- Files placed on SFTP, cloud storage, shared bucket
- Scheduled jobs read, validate, transform files
- Errors logged to error bucket or audit trail



Custom Protocols: Key Ideas

- Built for special constraints (performance, embedded systems, IoT)
- Uses bespoke binary framing, compression, or encoding
- Allows extreme optimization or special hardware interaction



Balwinder Sodhi

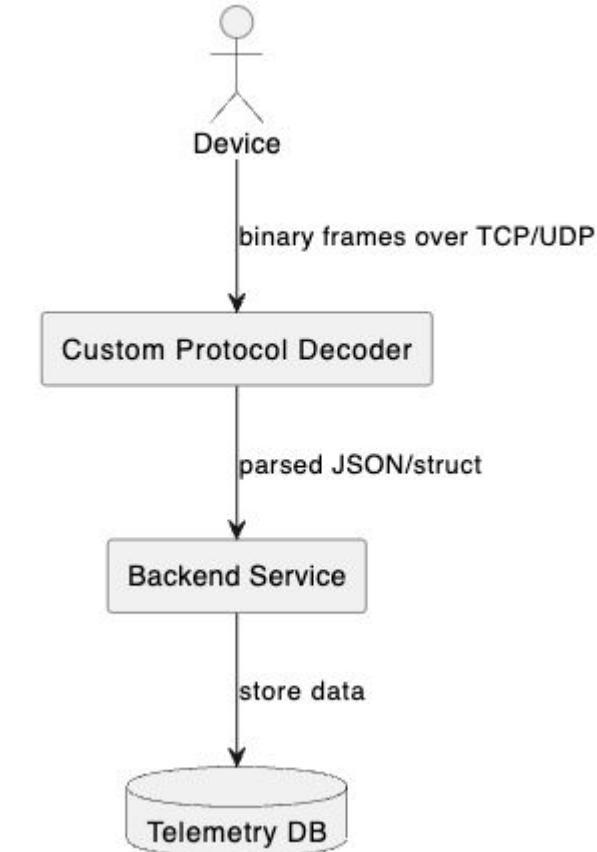
Custom Protocols: When to Use

- Real-time embedded systems
- Low-latency trading systems
- IoT constrained devices
- When standard protocols are too slow or heavy

Balwinder Sodhi

How Typically Implemented

- Custom serialization formats (CBOR, Cap'n Proto, FlatBuffers)
- Own framing rules
- Runs over TCP/UDP/QUIC
- Often includes custom error codes



Balwinder Sodhi

Serverless APIs: Key Ideas

- Functions-as-a-Service invoked by HTTP, events, or schedules
- Zero server management
- Automatic scaling
- Pay-per-invocation cost model



When to Use

MASER
DEVELOPMENT

A Guide For Early Career Engineers

- Bursty workloads
- Lightweight backend logic
- Event-driven automation
- MVPs or small teams without DevOps overhead



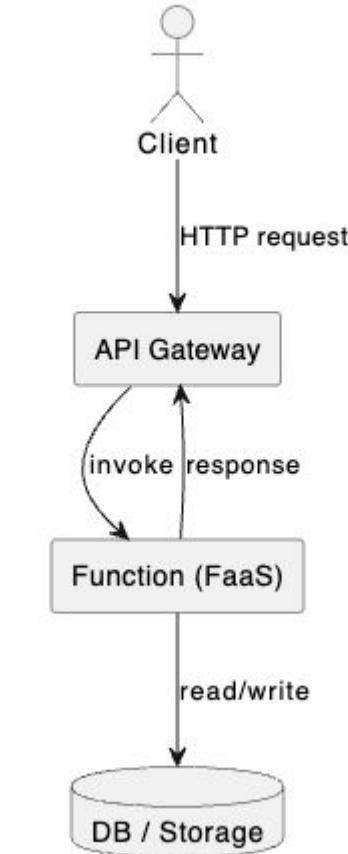
Balwinder Sodhi

How Typically Implemented

- API Gateway provides the HTTP endpoint
- Function runs business logic
- S3/DB/Queue integrations via triggers
- Authentication integrated (Cognito/OIDC)



Balwinder Sodhi



How to Choose an API Style

- REST → simple CRUD, wide client compatibility.
- GraphQL → flexible client-side data needs.
- gRPC → low-latency service-to-service calls.
- WebSockets → real-time updates.
- SOAP → strict enterprise contract governance.
- File-based → batch integrations.
- Serverless → bursty or event-driven workloads.

Balwinder Sodhi

Overall Summary

- Architectural design is about selecting a structure aligned with business, technical, and operational constraints.
- No single architecture fits all cases; trade-offs matter.
- Styles differ in modularity, scalability, cost, governance, and complexity.
- Cross-cutting concerns must be addressed regardless of chosen architecture.
- API style selection depends on usage patterns, consumers, and performance constraints.

Balwinder Sodhi