

# LLM-powered and Agentic Applications

Balwinder Sodhi

# What This Chapter Covers

- What counts as an LLM-powered application
- When agentic patterns are justified
- Architecture of agentic systems
- Core layers and interactions
- Architectural foundations and patterns
- How to think about trade-offs
- Practical guidance

Balwinder Sodhi

# What is an LLM-powered Application?

- A software system where LLM output meaningfully influences system behaviour.
- LLM is not the whole system; it's an embedded computational capability.
- Typical use:
  - Text generation
  - Semantic retrieval
  - Classification / extraction
  - Reasoning and planning at small scale

Balwinder Sodhi

# Anatomy of an LLM-Powered Application

- **App logic remains deterministic:** LLM augments it, not replaces it.
- **Common components:**
  - Prompting & templates
  - Vector search / RAG
  - Lightweight post-processing
  - Guardrails / validation
- **Interaction pattern:** **request → LLM → validate → integrate into workflow**

Balwinder Sodhi

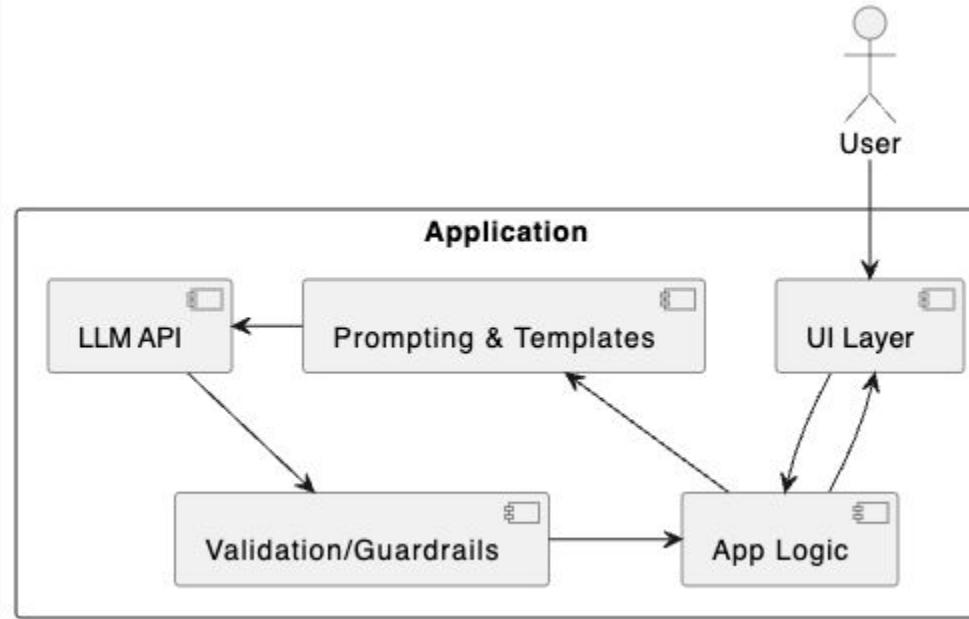
# High-Level Structure of an LLM-Powered App

MASTERING SOFTWARE  
DEVELOPMENT

A Guide For Early Career Engineers



Balwinder Sodhi



# Should You Build an Agentic Application or a Plain LLM-powered One?

You likely do not need agentic architectures when:

- The task is short-lived, stateless
- Input/output types are simple and well-defined
- No multi-step planning needed
- Errors are tolerable with retries
- Business logic is deterministic and bounded

Examples: summarization, classification, RAG chat, form filling, rewriting.

Balwinder Sodhi

# When You Should Consider Agentic Applications

Agentic architectures help when:

- You need multi-step workflows
- Actions depend on dynamic observation
- The system must self-correct
- External tools must be called in a feedback loop
- Long-running activities (minutes to hours)
- Goals cannot be achieved in one LLM call
- State is essential (context, memory, planning)

Balwinder Sodhi

Examples: autonomous research agents, automated operations (ticket triage → fix), data pipeline agents, orchestration workflows.

# Comparing the Two

Dimension	LLM-Powered App	Agentic App
State	Mostly stateless	Maintains evolving state
Execution	Single request-response	Multi-step loop
Control	Developer-driven	Goal-driven with agent feedback
Failure handling	Retry	Plan revision, tool fallback
Complexity	Low	High
Best for	Assistive tasks	Autonomous workflows

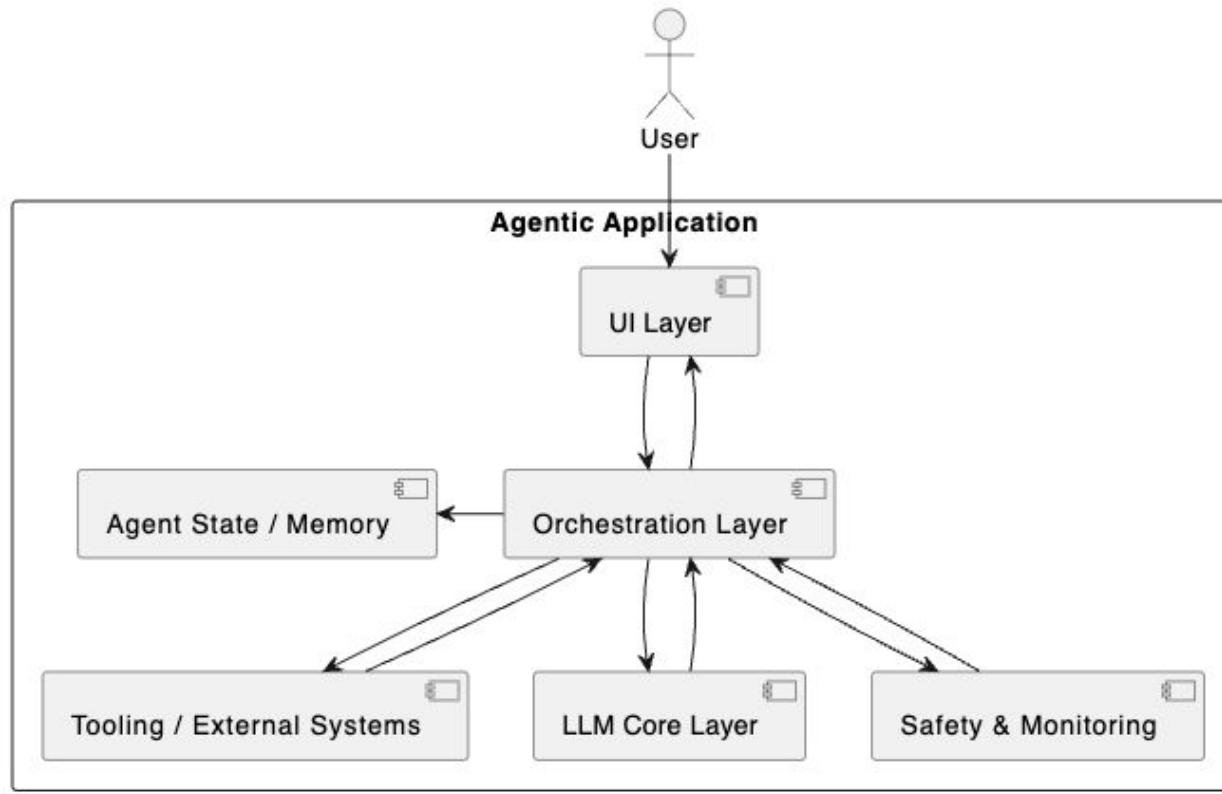
# Architecture of an Agentic Application

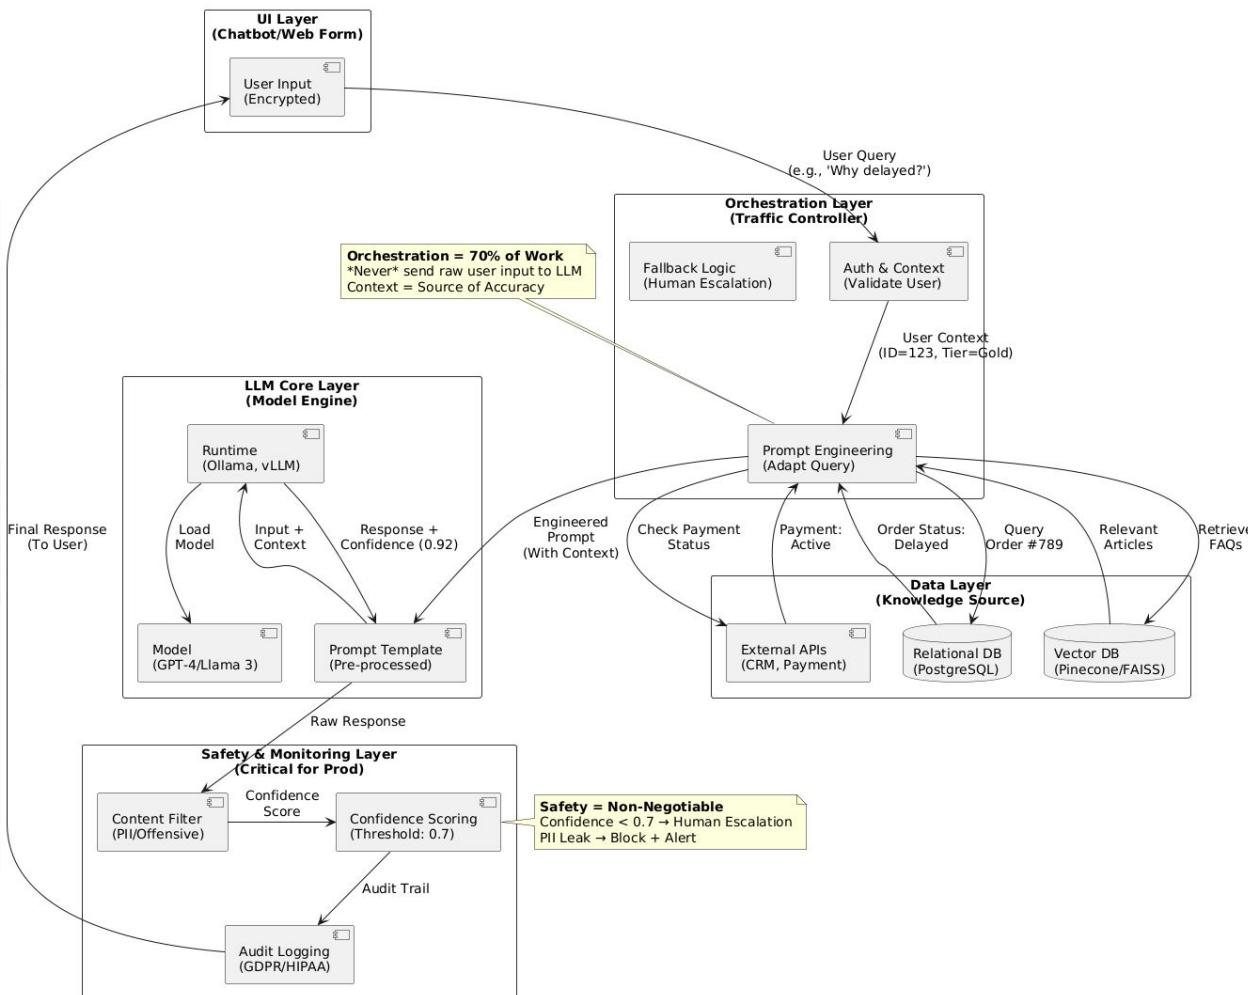
At a high-level it has five major layers:

1. UI Layer
2. Orchestration Layer
3. Data Layer
4. LLM Core Layer
5. Safety & Monitoring Layer

Balwinder Sodhi

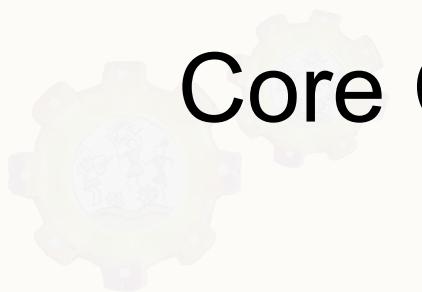
# High-level Architecture of an Agentic Application





MASTERING SOFTWARE  
DEVELOPMENT

A Guide For Early Career Engineers



# Core Components and Interactions

Balwinder Sodhi

# UI Layer

- Could be web, mobile, CLI, API endpoint
- Responsibilities:
  - Collect intents
  - Display results
  - Maintain interaction history
  - Provide grounding (context injection)
- Avoid embedding business logic here.

Balwinder Sodhi

# Orchestration Layer

- The heart of agentic systems.
- Responsibilities:
  - Loop management
  - State management
  - Selecting LLM prompts / strategies
  - Invoking tools
  - Handling fallbacks
  - Termination conditions
  - Logging, metrics, tracing
- This is where deterministic code constrains non-deterministic model behaviour.

# Data Layer

- Contains:
  - Long-term memory (vector DB, relational DB)
  - Short-term scratchpads
  - Context stores
  - Knowledge bases
  - Cached computations
- Functions:
  - Retrieval
  - Storing agent state
  - Persisting workflow history
  - Guarding against hallucinations through retrieved ground truth

# LLM Core Layer

Mainly includes:

- Model endpoints (OpenAI, Anthropic, local models)
- Prompt templates
- Sampling strategies (temperature, top-p)
- Tool-use capabilities
- System instructions
- Structured output handling (JSON schemas)

Balwinder Sodhi

# Safety and Monitoring Layer

MATHEW  
Covers: SOFTWARE  
DEVELOPMENT

- Input/output guardrails
- Toxicity / jailbreak detection
- Role-based access controls for agent actions
- Observability: traces, metrics, logs
- Replay systems
- Circuit breakers for runaway loops
- Quotas & budget monitoring

MASTERING SOFTWARE  
DEVELOPMENT

A Guide For Early Career Engineers



# Architectural Foundations of LLM-Powered Software

Balwinder Sodhi

# Key Architectural Challenges

- **Non-determinism:** LLM outputs vary
- **Steering:** aligning model output with business needs
- **Bounded behaviour:** preventing runaway loops or harmful actions
- **Debuggability:** difficult due to probabilistic reasoning
- **Performance optimization:** cost vs responsiveness
- **State management:** context growth, memory pruning
- **Safety:** restricting tool use and actions

Balwinder Sodhi

# Core Idea

MASTERING SOFTWARE  
DEVELOPMENT

*LLM-powered architectures must add deterministic structure around a non-deterministic core.*

This shapes most patterns in modern AI systems.

Balwinder Sodhi

# The Agentic Architecture Pillars: #1 Structure

The agent must operate inside well-defined boundaries:

- Allowed actions
- Goal templates
- Planning constraints
- Valid schemas
- Safety limits

Balwinder Sodhi

# Pillar 2: Grounding

ML  
DATA  
DEVELOPMENT  
A Q de F

LLMs must operate on retrieved and validated context:

- Documents
- Databases
- Tool results
- Execution logs

Improves correctness and reduces hallucinations.

Balwinder Sodhi

# Pillar 3: Observability

Agents must produce:

- Traces (per-step)
- Logs
- Model inputs/outputs
- Replayable workflows
- Error snapshots

Balwinder Sodhi

# Pillar 4: Deterministic Control

The orchestrator must:

- Control loops
- Enforce state transitions
- Validate output
- Govern safety and tool usage

Deterministic code contains the model.

Balwinder Sodhi

MASTERING SOFTWARE  
DEVELOPMENT

A Guide For Early Career Engineers



# Architectural Patterns: Making Trade-offs Explicit

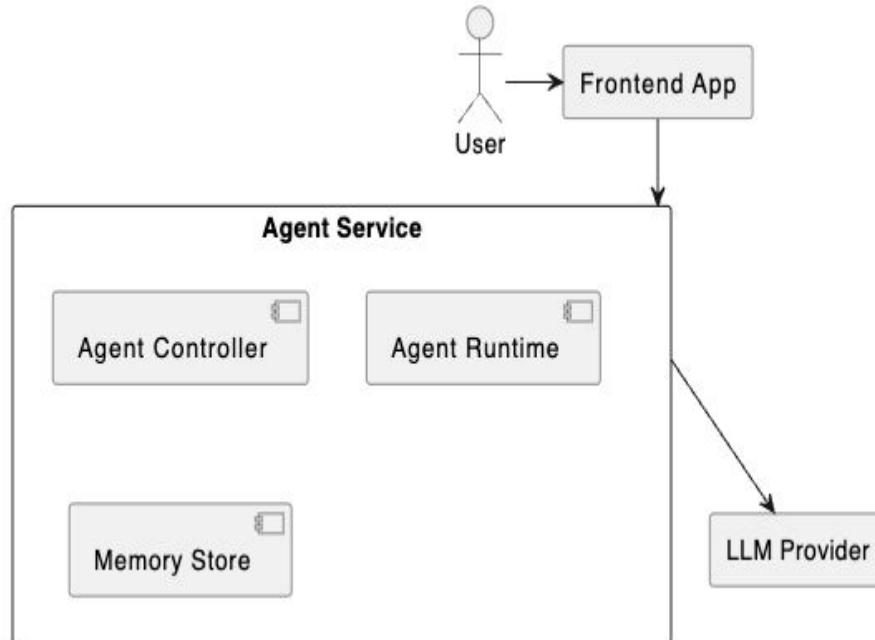
Balwinder Sodhi

# Agent as a First-Class Service

MASTERING SOFTWARE  
DEVELOPMENT

Treat the agent like any other backend service:

- Has an API
- Maintains state
- Exposes capabilities
- Instances can be scaled horizontally
- Decouples agent execution from the UI



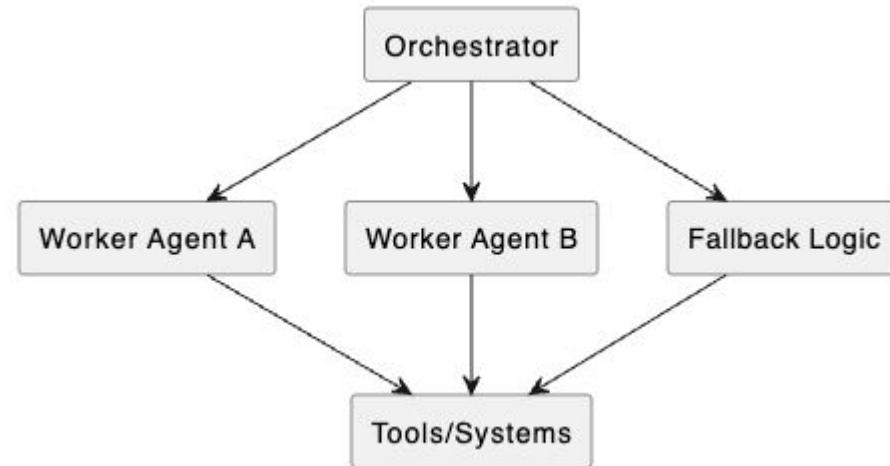
# Orchestrator–Worker With Explicit Fallbacks

**Key idea:** Use an orchestrator agent that delegates tasks to worker agents or tools, with fallback logic baked in.

A Guide For Early Career Engineers

Benefits:

- Robustness under uncertainty
- Clear responsibility boundaries
- Modular task decomposition

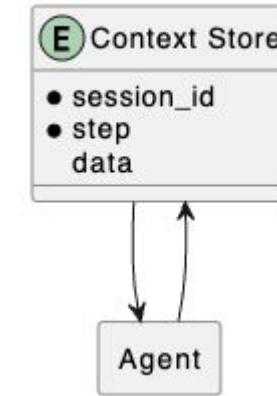


Balwinder Sodhi

# Context as a First-Class Data Type

Key idea: Context must be explicitly managed:

- Passed to every agent step
- Versioned
- Composable
- Persisted for replay
- Stored in structured formats



Prevents “lost context” errors and uncontrollable behaviour.

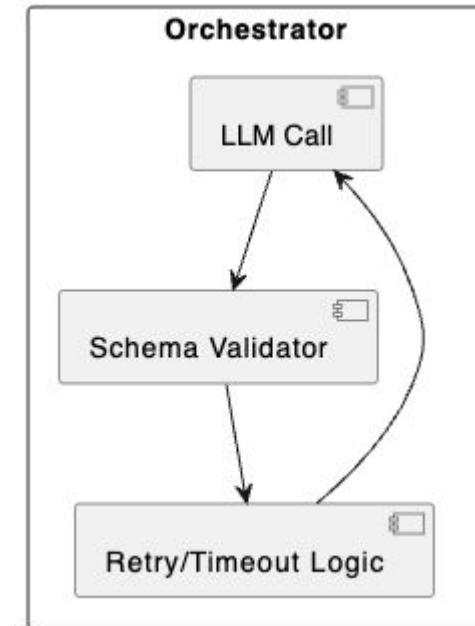
Balwinder Sodhi

# Deterministic Boundaries for Non-Deterministic Systems

All LLM outputs are validated:

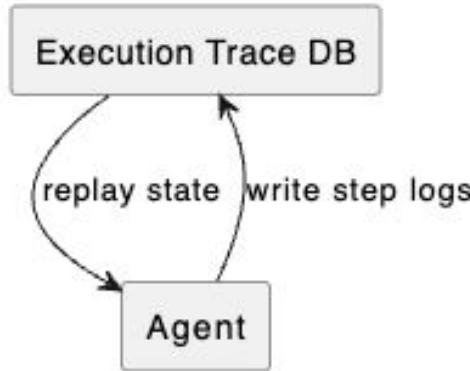
- JSON schema enforcement
- Enum constraints
- Tool selection constraints
- Step timeouts
- Retry policies
- Max depth, max steps

Balwinder Singh  
Deterministic boundaries create safe operational envelopes.



# Observability and Replayability

- Stores:
  - Prompts
  - Inputs
  - Outputs
  - Tool invocation logs
  - Execution traces
- Used for:
  - Debugging
  - Governance
  - Safety audits
  - Improvement cycles

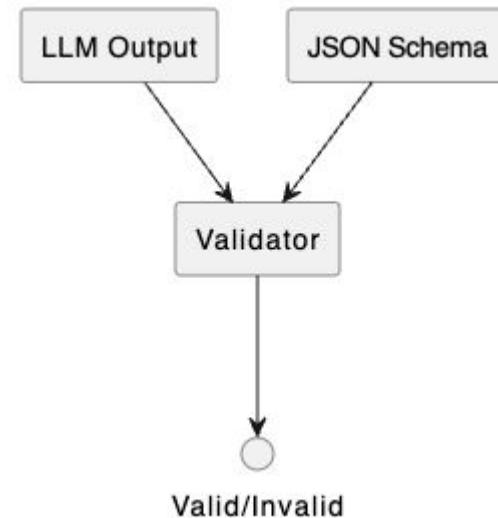


# Schema-Governed Contracts

LLM output must conform to explicit schemas:

- JSON schema
- DSLs
- Strict enums
- Bounded numeric ranges
- Plan schema (steps, actions, dependencies)

Helps constrain non-deterministic text generation.



Balwinder Sodhi

# Key Takeaways

- Not all LLM usage requires agents
- Agentic apps introduce complexity—use them wisely
- Architecture must constrain LLM unpredictability
- Orchestration is where most engineering value lies
- Observability and safety are essential
- Context and schemas reduce hallucinations
- Patterns help structure systems under uncertainty

Balwinder Sodhi