

Database Design and Management

Balwinder Sodhi

Why Database Design Shapes System Quality

- Databases persist the most long-lived state of the system.
- Structural mistakes accumulate compound interest:
 - Poor joins → systemic latency
 - Unbounded growth tables → storage & performance bottlenecks
 - Incorrect relationships → inconsistent user flows
 - Wrong normalization → slow writes or slow reads

A Guide For Early Career Engineers



Balwinder Sodhi

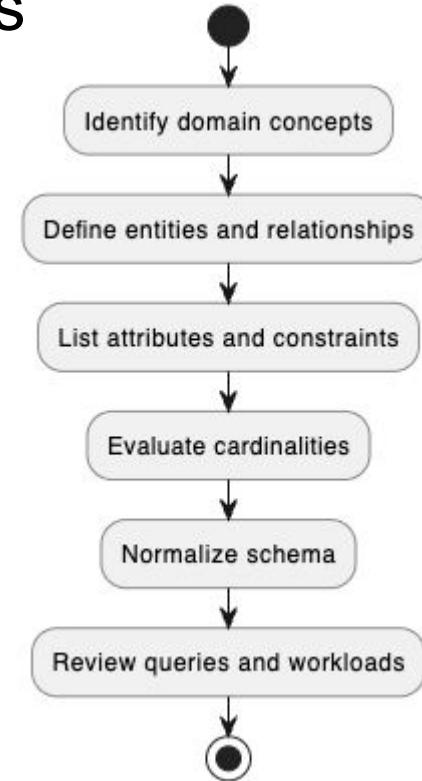
Data Modeling: Sketching Before Building

- Domain modeling identifies concepts, rules, boundaries.
- Data modeling transforms domain concepts into entities, attributes, constraints.
- Good modeling avoids premature optimization:
 - Keep schema expressive but simple.
 - Avoid leaking UI and API concerns into schema.

Balwinder Sodhi

Example Data Modeling Thought Process

- **Domain:** Online Bookstore
- **Identify entities:** Users, Books, Orders, Payments, Inventory, Reviews
- **Identify rules:**
 - A user can have multiple orders
 - A book can have multiple reviews
 - An order can contain many books
- **Identify constraints:**
 - Stock quantity cannot go negative
 - Order total must equal sum of items



A Simple Online Bookstore

Core Entities

MASTERING SOFTWARE DEVELOPMENT

A Guide For Early Career Engineers

- Users
- Books
- Orders
- Payments
- Reviews
- Inventory

Typical Workflows

- User browses books → search/indexing considerations
- Creates an order → transactions, locking, consistency
- Reduces stock → constraints & triggers
- Writes reviews → preventing duplicates, spam protection

Balwinder Sodhi

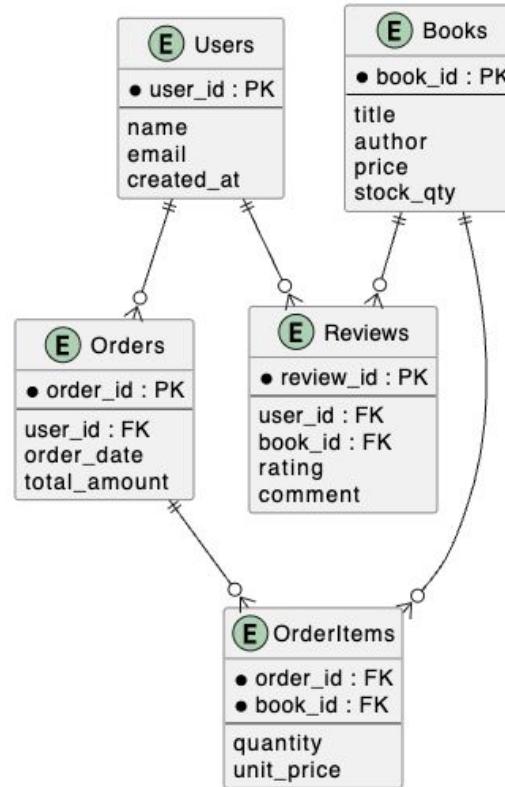
Online Bookstore Entities

MASTERING SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers



Balwinder Sodhi



Types of Relationships

- One-to-One (1:1)
- One-to-Many (1:M)
- Many-to-Many (M:N)



Balwinder Sodhi

One-to-One (1:1)

MASER SHARE

DEVELOPMENT

A Guide For Early Career Engineers

When to Use

- Each row in Table A corresponds to at most one row in Table B.
- Entities conceptually distinct but tightly coupled.
- Sensitive or rarely accessed data separated for security/performance.
- Very wide tables broken up into logical modules.

Common Real-World Uses

- Users → UserProfiles
- Orders → OrderPayments
- Employees → EmployeeConfidentialInfo

One-to-One (1:1)

MASER DEVELOPMENT

- Implementation Patterns

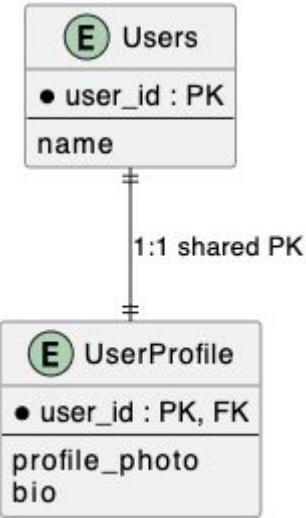
- Shared Primary Key (strongest constraint)

A Guide For Early Career Engineers | <https://msd-book.github.io/>
B's primary key is also a foreign key to A.

- Foreign key in one direction (soft 1:1)
B has a unique constraint on FK to A.

- Trade-offs

- Shared PKs strictly enforce cardinality but add migration complexity.
- Soft 1:1s offer flexibility but risk violating true uniqueness unless constraints in place.



Balwinder Sodhi

One-to-Many (1:N)

- Definition: One row in Table A connects to zero or more rows in Table B.
- It fits perfectly when you have a “most natural” relationship in business domains:
 - A customer → multiple orders
 - A book → many reviews
 - A device → many sensor readings
- Implementation
 - B contains a foreign key pointing to A.

Balwinder Sodhi

One-to-Many (1:N)

MAS
TER
DEVELOPMENT

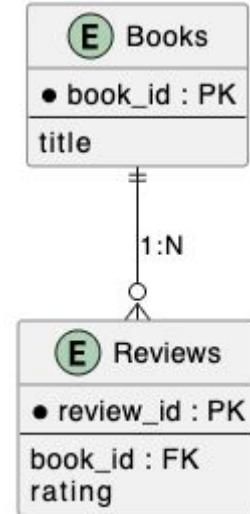
A Guide For Early Career Engineers

- Performance Considerations

- FK lookup on B→A is fast if indexed.
- Querying all children for parent is straightforward.
- Cascading deletes require careful reasoning in production workloads.

- When It Becomes a Problem

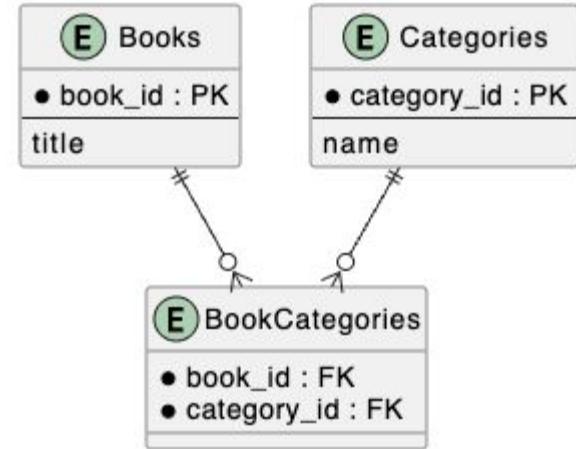
- Child table grows unbounded (e.g., logs, metrics, events).
- FK introduces locking contention on writes.



Balwinder Sodhi

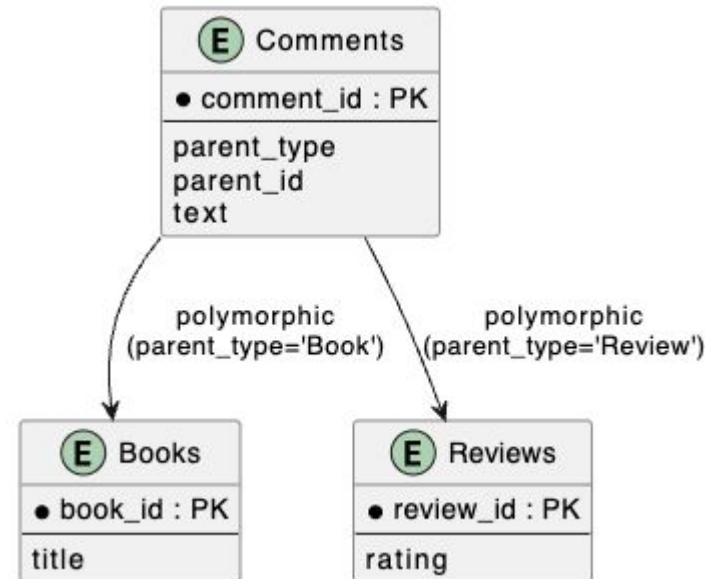
Many-to-Many (M:N)

- Definition: A row in A relates to multiple rows in B, and vice versa.
- Always implemented through a junction table (a.k.a. join table or association table).
- Examples:
 - Students ↔ Courses
 - Books ↔ Categories
 - Users ↔ Roles
- Advanced Practical Considerations
 - The join table can grow extremely large → impacts join performance.
 - Queries often require composite indexes (e.g., (A_id, B_id)).
 - Cascades on many-to-many relationships must be thought through carefully.



Polymorphic Relationships

- One table references multiple other tables via a “type” + “id” pair.
- Pros
 - Flexible
 - Reduces table explosion for similar concepts
- Cons
 - No FK constraints possible → referential integrity risk
 - Harder to query efficiently
 - Requires type-safe checks in app logic



Heuristics for Deciding Relationship Types

Ask these questions:

- Is this object dependent on another object to exist?
 - Yes → likely 1:N or 1:1 (composition)
 - No → maybe M:N or separate domain
- How does the domain describe the relationship?

Natural language helps:

- “A user has many orders”
- “An order contains multiple items”
- “A product belongs to a category”
- These map 1:1 to schema choices.

Heuristics for Deciding Relationship Types

Ask these questions:

- What is the anticipated query pattern?
 - If you always query books with authors → embedding or denormalizing may help.
 - If you frequently need category filters → M:N join table with proper indexing.
- What is the data growth pattern?
 - Logs/events → enormous 1:N → consider vertical partitioning
 - Many-to-many between large domains → massive join tables → consider caching or search engine indexing

Balwinder Sodhi

Hard Referential Integrity (Database-Enforced)

- Using foreign key constraints.

- Pros:

A Guide For Early Career Engineers

- Guaranteed correctness
- Catches bugs early
- Maintains domain consistency
- Enables cascading rules

- Cons:

- Slows write-heavy systems
- Causes locking issues
- Hard to bulk-update or run large migrations
- Multi-region DB setups complicate FK guarantees

Soft Referential Integrity (App-Enforced)

- Application code checks relationships; database stores plain IDs without FK constraints.
- Pros:
 - Faster writes
 - Easier sharding
 - Lower lock contention
 - Ideal for event-driven and append-only systems
- Cons:
 - Requires robust tests + observability
 - Risk of broken references
 - Developers must build cleanup tasks & orphan detection
- When Used in Practice
 - Twitter, Meta, etc. with massive write loads
 - Systems using Kafka-like async flows
 - Federated architectures with independent services

Cardinality Impacts Query Efficiency

- 1:1 → predictable joins, simplest queries
- 1:N → often OK, but watch child table size
- M:N → worst-case complexity increases dramatically
- Polymorphic → kills most optimizers

Balwinder Sodhi

Normalization: Reducing Redundancy and Improving Integrity

Why Normalize?

- Prevent duplicate/contradictory data
- Make updates consistent by default
- Ensure schema represents domain constraints
- Reduce bugs from hidden transitive dependencies

Balwinder Sodhi

Step-by-Step Normalization Example (Book + Author)

Imagine a CSV imported from legacy system (unnormalized data):

StudentID	StudentName	Courses	Instructor
1	Arun	Math, Physics	Prof. Singh
2	Barun	Math	Prof. Singh
3	Charlie	Physics, Chemistry	Prof. Singh, Prof. Nath

- Issues:
 - Multi-valued Courses field (violates 1NF)
 - Instructor data repeated for every enrolment
 - Instructor depends on Course (transitive dep)

1NF: Ensure Atomic Columns

To transform the table to First Normal Form (1NF) we ensure that the columns contain only atomic values as in table below:

A Guide For Early Career Engineers

StudentID	StudentName	Course	Instructor
1	Arun	Math	Prof. Singh
1	Arun	Physics	Prof. Singh
2	Barun	Math	Prof. Singh
3	Charlie	Physics	Prof. Singh
3	Charlie	Chemistry	Prof. Nath

Second Normal Form (2NF)

- We make sure it is in 1NF and that the non-key attributes must depend on the whole primary key.
- Current primary key is (StudentID, Course).
- To fix the problem, we decompose the table into two:

StudentID	Student Name
1	Arun
2	Barun
3	Charlie

StudentID	Course	Instructor
1	Math	Prof. Singh
1	Physics	Prof. Singh
2	Math	Prof. Singh
3	Physics	Prof. Singh
3	Chemistry	Prof. Nath

Third Normal Form (3NF)

- The problem with the above (2NF) tables is that Instructor depends on Course, not directly on the key (StudentID, Course).
- To fix it, we have to have the tables in 2NF and no transitive dependencies (i.e., non-key attributes shouldn't depend on other non-key attributes) in the table columns.
- Enforcing this idea gives us the Third Normal Form (3NF) as shown in the decomposed tables below.

Balwinder Sodhi

Third Normal Form (3NF)

StudentID	Student Name
1	Arun
2	Barun
3	Charlie

StudentID	Course
1	Math
1	Physics
2	Math
3	Physics
3	Chemistry

Course	Instructor
Math	Prof. Singh
Physics	Prof. Singh
Chemistry	Prof. Nath

A table is in 1NF when

MASTERING SOFTWARE
DEVELOPMENT

- All values are atomic (no repeating groups, arrays, or composite values).
- Each record is unique (no duplicate rows).
- Each column contains values of a single data type.
- The order of rows and columns does not matter.

Balwinder Sodhi

A table is in 2NF when

- It is already in 1NF.
- Every non-key attribute is fully functionally dependent on the entire primary key.
 - This applies only when the primary key is composite.
 - No partial dependencies (i.e., a non-key attribute depending on only part of a composite key).

Balwinder Sodhi

A table is in 3NF when

- It is already in 2NF.
- There are no transitive dependencies:
 - No non-key attribute depends on another non-key attribute.
 - All non-key attributes must depend directly and only on the primary key.

MAS

ERIN

DEVELOPMENT

A Guide For Early Career Engineers



Balwinder Sodhi

MASTERING SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers

Triggers and Stored Procedures/Functions

Balwinder Sodhi

What Triggers Are

- Database-side automation reacting to INSERT/UPDATE/DELETE.

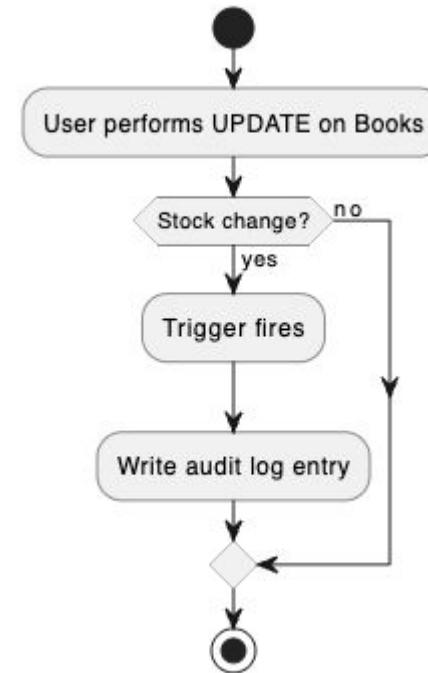
● Use cases:

A Guide For Early Career Engineers

- Audit trails
- Maintaining derived values
- Enforcing complex constraints
- Cascading behavior

● Pitfalls:

- Harder to debug
- Hidden logic → surprising for developers
- Performance impact if overused
- Avoid business logic in triggers



Stored Procedures & Functions: Reusable Database Code

● When to Use Them

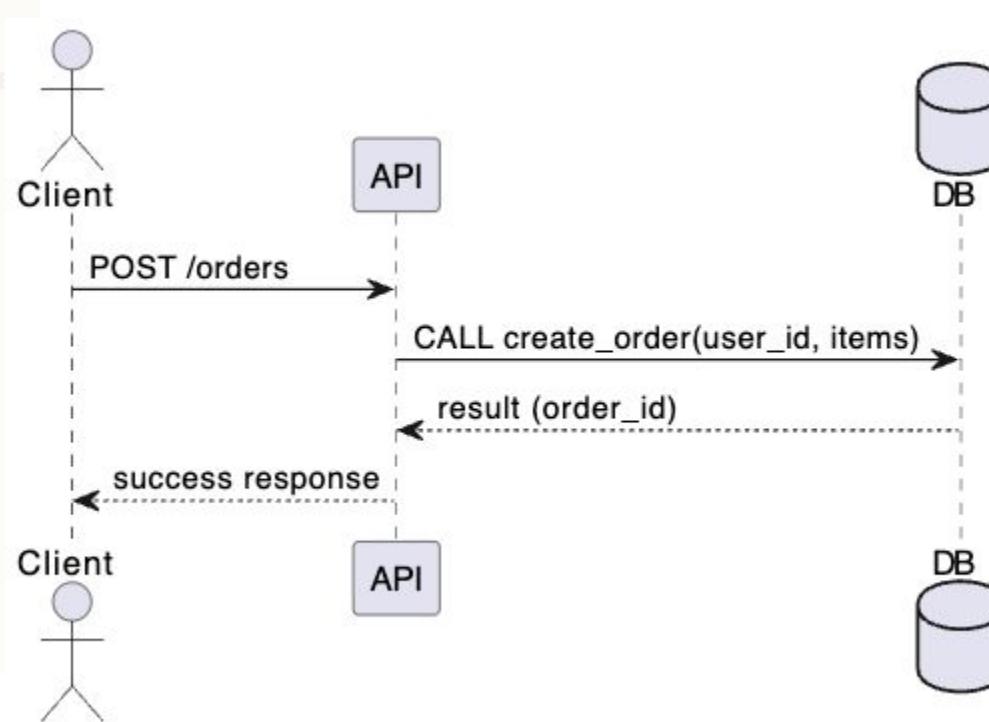
- Encapsulate multi-step operations (e.g., creating an order).
- Improve performance by reducing network round trips.
- Enforce security boundaries (role-based access).
- Simplify batch updates / heavy data transforms.

● Best Practices

- Keep them small and purpose-driven.
- Version-control them like application code.
- Avoid embedding business logic that should live in services.
- Prefer idempotent operations where possible.

Balwinder Sodhi

Stored Procedure Call Example



Advanced Trigger Use Cases

- Materialized view maintenance
- Change Data Capture (CDC) pipelines
- Audit history + temporal tables
- Cache invalidation (Redis, CDN)

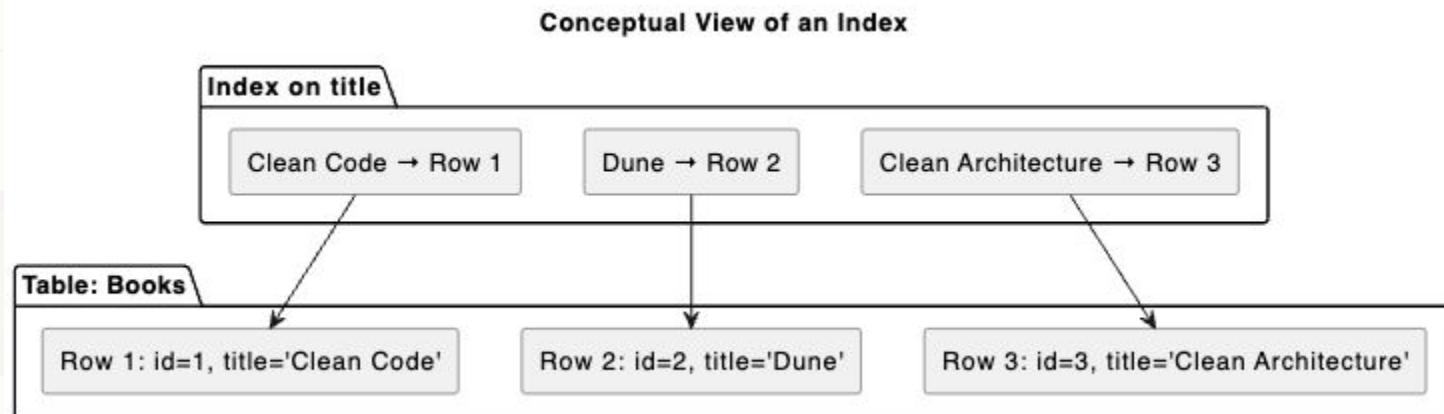
Avoid in Triggers:

- Calling external APIs (timeouts kill transactions)
- Heavy computation
- Cross-row loops (N^2 behavior)

What Is an Index?

- An index is a data structure that helps the database find rows faster, just like an index in the back of a book:
 - Without an index → the database reads the whole table (full scan).
 - With an index → the database jumps directly to relevant rows.
 - Indexes trade faster reads for:
 - slower writes (because index updates are needed)
 - extra storage
- Real-World Analogy
 - Think of a phone directory: names sorted alphabetically.
To find “Alice”, you don’t read every page — you jump directly to the A section.

Conceptual view



Balwinder Sodhi

Why Indexes Make Queries Faster

- Indexes allow the database to avoid scanning irrelevant rows.

Two Common Query Paths:

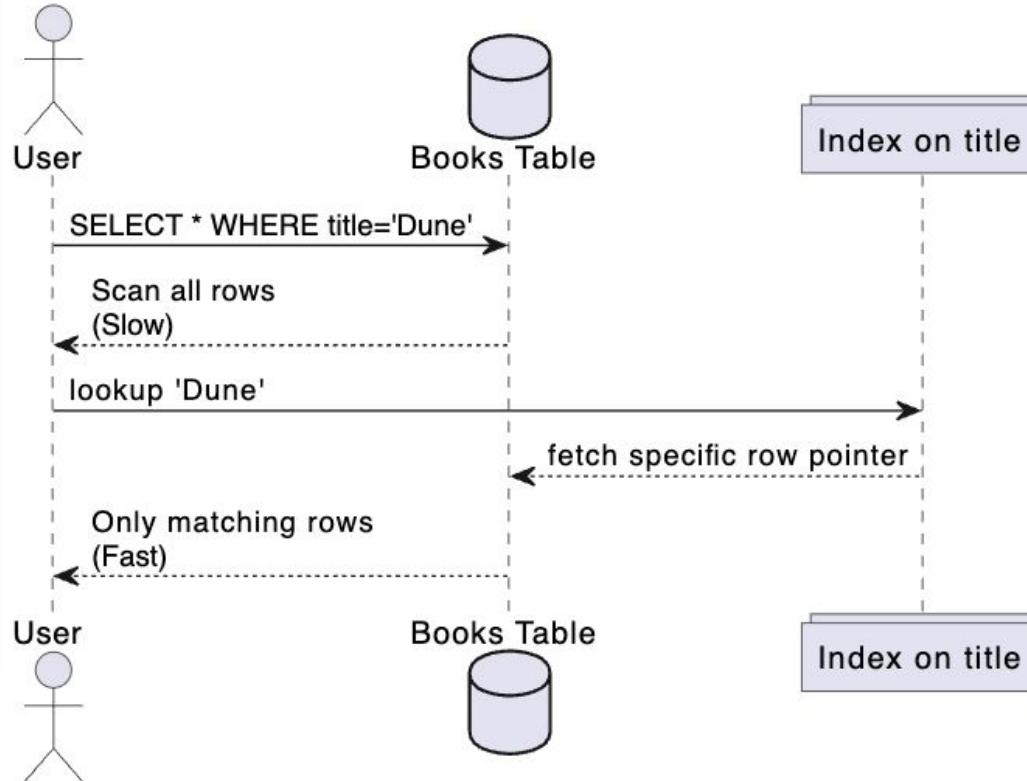
- Without Index (Full Table Scan)
 - DB reads every row → checks the condition
 - Slow on large tables ($O(n)$)
- With Index (Index Seek)
 - DB uses a sorted/searchable structure
 - Quickly locates matching values ($O(\log n)$)
 - Fetches only required rows

When Indexes Help

- Searching by a column (e.g., ISBN, email, username)
- Sorting (ORDER BY)
- Filtering (WHERE)
- Joins on foreign keys
- Partial match text search (GIN/GiST)

Balwinder Sodhi

Full Scan vs Index Seek



Commonly Used Types of Indexes

- B-Tree indexes
 - Default in most RDBMS
 - Great for ranges, sorting, equality
- Hash indexes
 - Equality only
- GIN / GiST indexes
 - Document/JSONB search
 - Full-text search
 - Array membership queries
- Composite indexes
 - Multi-column indexing
 - Order matters
- Covering indexes
 - Index contains all fields needed → no table read

Transactions & Isolation Levels

ACID Principles

- Atomicity — all or nothing
- Consistency — valid state transitions
- Isolation — concurrent operations behave independently
- Durability — data survives crashes

Balwinder Sodhi

Transaction Isolation Levels

Level	Prevents	Allows
Read Uncommitted	nothing	dirty reads
Read Committed	dirty reads	non-repeatable reads
Repeatable Read	dirty + non-repeatable reads	phantom reads (varies)
Serializable	all anomalies	slowest, uses locks/predicates

Anomaly Examples

• Dirty Read

- Transaction A updates row
- Transaction B reads uncommitted data
- A rolls back → B saw ghost data

• Non-Repeatable Read

- Transaction A reads row
- Transaction B updates same row
- A reads again → gets different result

• Phantom Read

- Transaction A queries range
- Transaction B inserts new matching rows
- A re-queries → sees extra rows

MASTERING SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers



Balwinder Sodhi

Performance Tuning

Where Performance Problems Usually Come From

- Missing indexes
- Bad cardinality estimates
- Too many joins
- Unbounded scans
- Large transactional locks
- Poor schema normalization
- Hotspot rows (e.g., "global counters")

Balwinder Sodhi

Query-Level Optimization

- Avoid SELECT *
- Filter early (in SQL, not app)
- Limit result sets
- Avoid unnecessary ORDER BY or DISTINCT
- Use materialized views for heavy analytic queries

Balwinder Sodhi

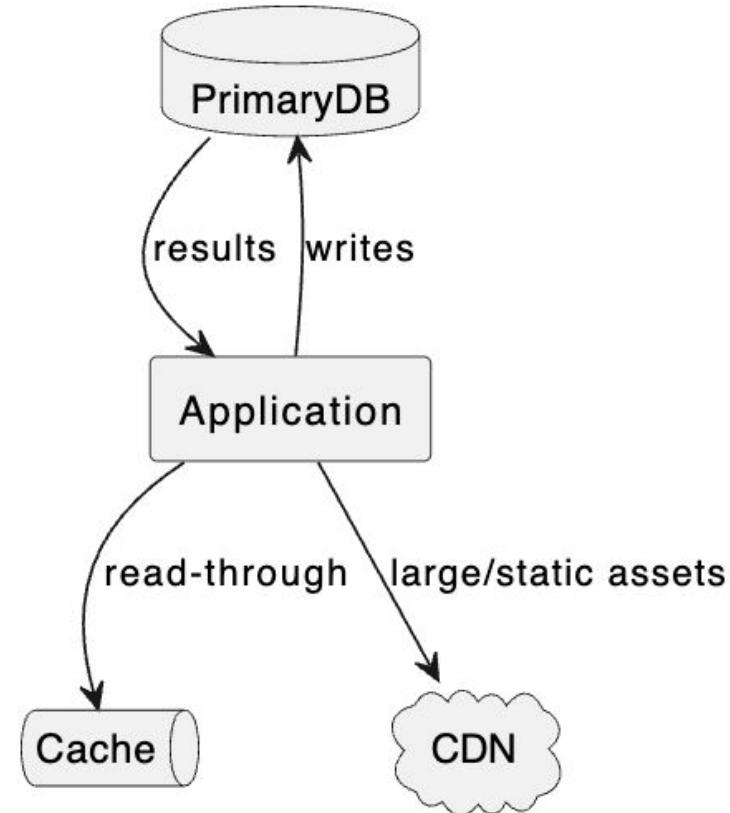
Schema-Level Optimization

- Use proper data types
- Partition large tables
- Avoid storing large blobs inline
- Introduce summary/rollup tables
- Use foreign key indexes

Balwinder Sodhi

Application-Level Optimization

- Use caching layers (Redis)
- Batch writes
- Apply optimistic concurrency where possible
- Use read replicas for heavy read traffic



Balwinder Sodhi

MASTERING SOFTWARE
DEVELOPMENT

A Guide For Early Career Engineers

Beyond Relational: A Glimpse into NoSQL

Balwinder Sodhi

Why NoSQL Exists

MASER
DEVELOPMENT
A Guide For Early Career Engineers

- Scaling relational systems horizontally is difficult.
- Modern applications store semi-structured and large-volume data.
- NoSQL trades strict schema + ACID for flexibility + scalability.



Balwinder Sodhi

Major NoSQL Types

- Document Stores (MongoDB)
 - Great for nested data, variable schemas
 - Natural for JSON-based workflows
- Key-Value Stores (Redis)
 - Fast caching & ephemeral data
- Wide-Column Stores (Cassandra)
 - High write throughput
 - Tunable consistency
- Graph Databases (Neo4j)
 - Relationship-heavy domains (routing, recommendations)

Document vs Relational Example (Book example)

- Relational: Books table, Authors table, BookCategory join table.
- Document store equivalent:

A Guide For Early Career Engineers

```
{  
    "book_id": "B1",  
    "title": "Clean Code",  
    "authors": ["Robert C. Martin"],  
    "categories": ["software", "engineering"],  
    "reviews": [  
        { "user": "U1", "rating": 5, "comment": "Excellent." }  
    ]  
}
```

Balwinder Sodhi

Choice of SQL vs NoSQL

- Choose NoSQL when
 - Schema flexibility needed.
 - Large-scale reads/writes.
 - Data naturally hierarchical or graph-like.
 - Event ingestion, logs, time-series workloads.
- When Not To
 - Strong consistency & transactions crucial.
 - Complex multi-table relationships that benefit from joins.

Balwinder Sodhi

Key Takeaways

- Start with modeling → schema-first approach reduces long-term cost.
- Normalize first; denormalize carefully based on real workloads.
- Understand relationships deeply; they shape query performance.
- Use triggers & stored procedures thoughtfully.
- NoSQL offers flexibility and scale, but with different trade-offs.
- Database design is not static—schemas evolve with the product.

Balwinder Sodhi